

- CS190I Program Synthesis for the Masses -

# Lecture 10:

# MARS: Maximal Multi-Layer Specification Synthesis

Yanju Chen

Computer Science Department

University of California, Santa Barbara

Spring 2021

# Reminders

- R3 is due on Wed 4/28 (today!)
- HW2 is due on Mon 5/3 (next week)
  - Start early, ask if you encounter any confusions
- Proposal is due on Wed 5/5 (next week)
  - Tips for writing a proposal -> [here](#)

# HW2 Frequently Asked Questions

- Racket is not found from my terminal even after I installed it.
  - You'll need to add the path to the Racket binaries to your system environment in order to execute `racket` as well as `raco` (Racket's package management tool). You can refer to this "full setup" [here](#).
  - For macOS, your path file is *usually* `~/.bash_profile` (if using `bash` as your default shell) or `~/.zshenv` or `~/.zshrc` (if using `zsh` as your default shell); For Linux, your path file is *usually* `~/.bashrc` (if using `bash` as your default shell).
  - Note that Venti requires Racket 7.7 and Rosette 3.2. The most recent versions of Racket 8.0 and Rosette 4.0 are major ones that are not backward compatible. So make sure you follow the hw instructions to have the correct versions configured.
- There's an exception when I tried to run the `yul_parser`.
  - It's possible that you are not using Antlr 4.8 Python Target. You can easily get it by `pip`:
    - `pip install antlr4-python3-runtime==4.8`
    - Also make sure your Python version is 3.0+.
  - If you have other questions, come to the office hours or message me on Slack 😊.

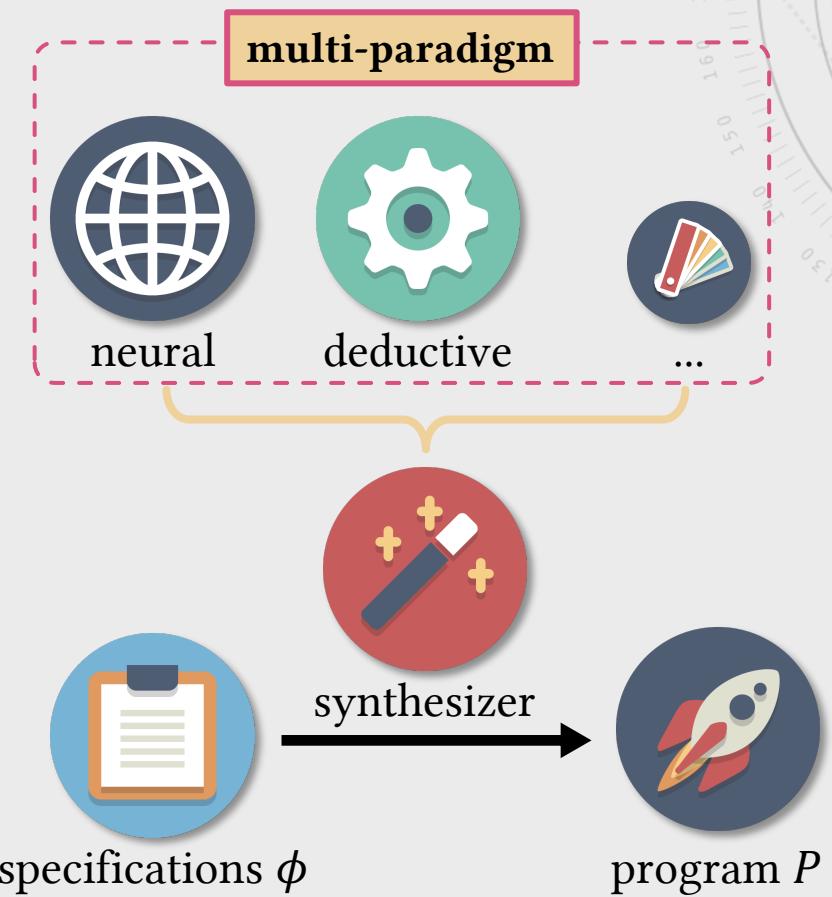
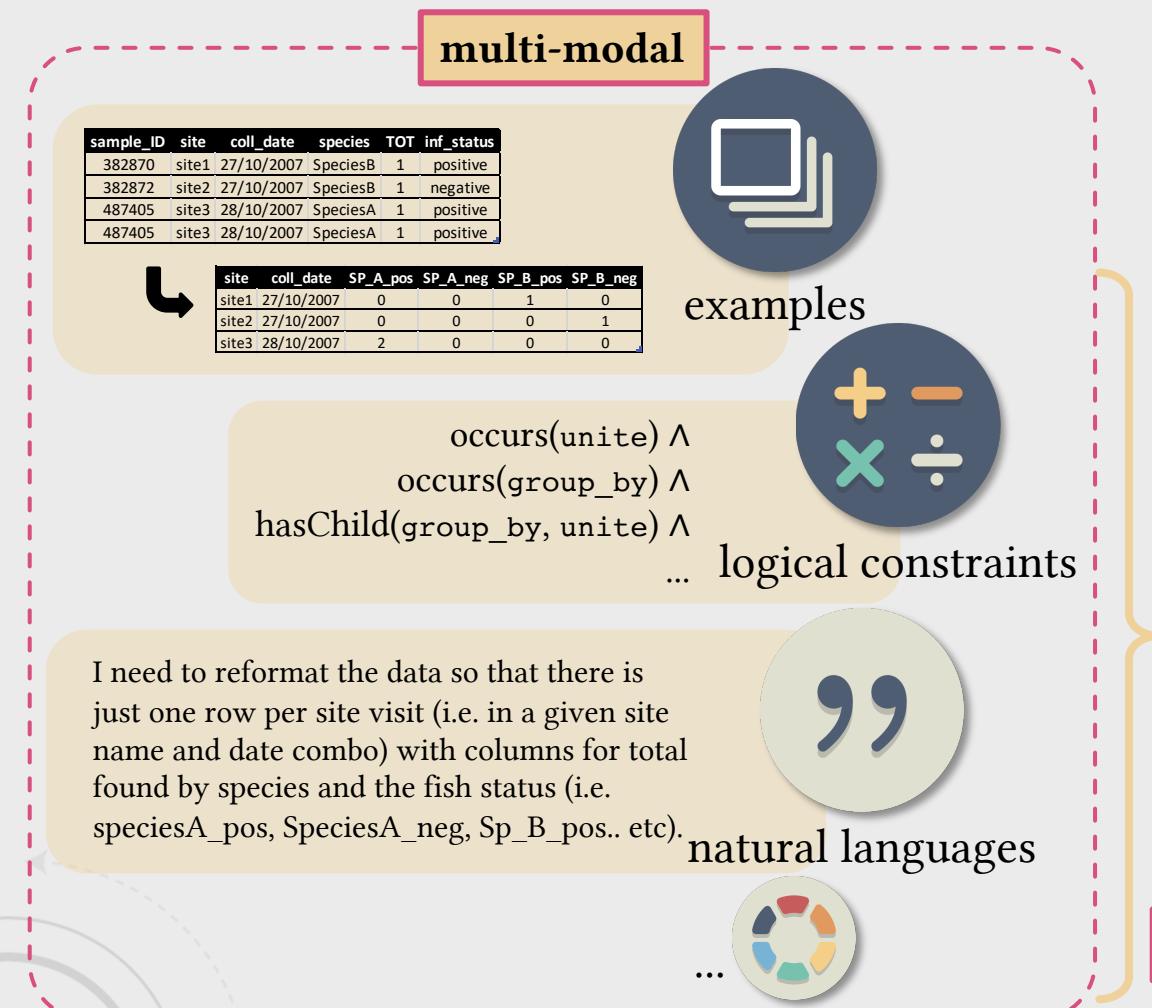
# Overview

- Program Synthesis in a Nutshell
- MARS: Encoding Multi-Layer Specifications
- Related Works & Conclusions

# Program Synthesis in a Nutshell

- Problem Formalization
  - Related Works
  - Program Synthesis with Machine Learning (I)
  - A Data Wrangling Example & DSL
  - NEO: A Brief Overview
- Observations & Motivations
    - Q1: Why logical reasoning?
    - Q2: Why machine learning?
    - Q3: Why bridging?

# Problem Formalization



**Find a program  $P$  that satisfies all the specifications  $\phi$ .**

# Program Synthesis in a Nutshell

## Related Works

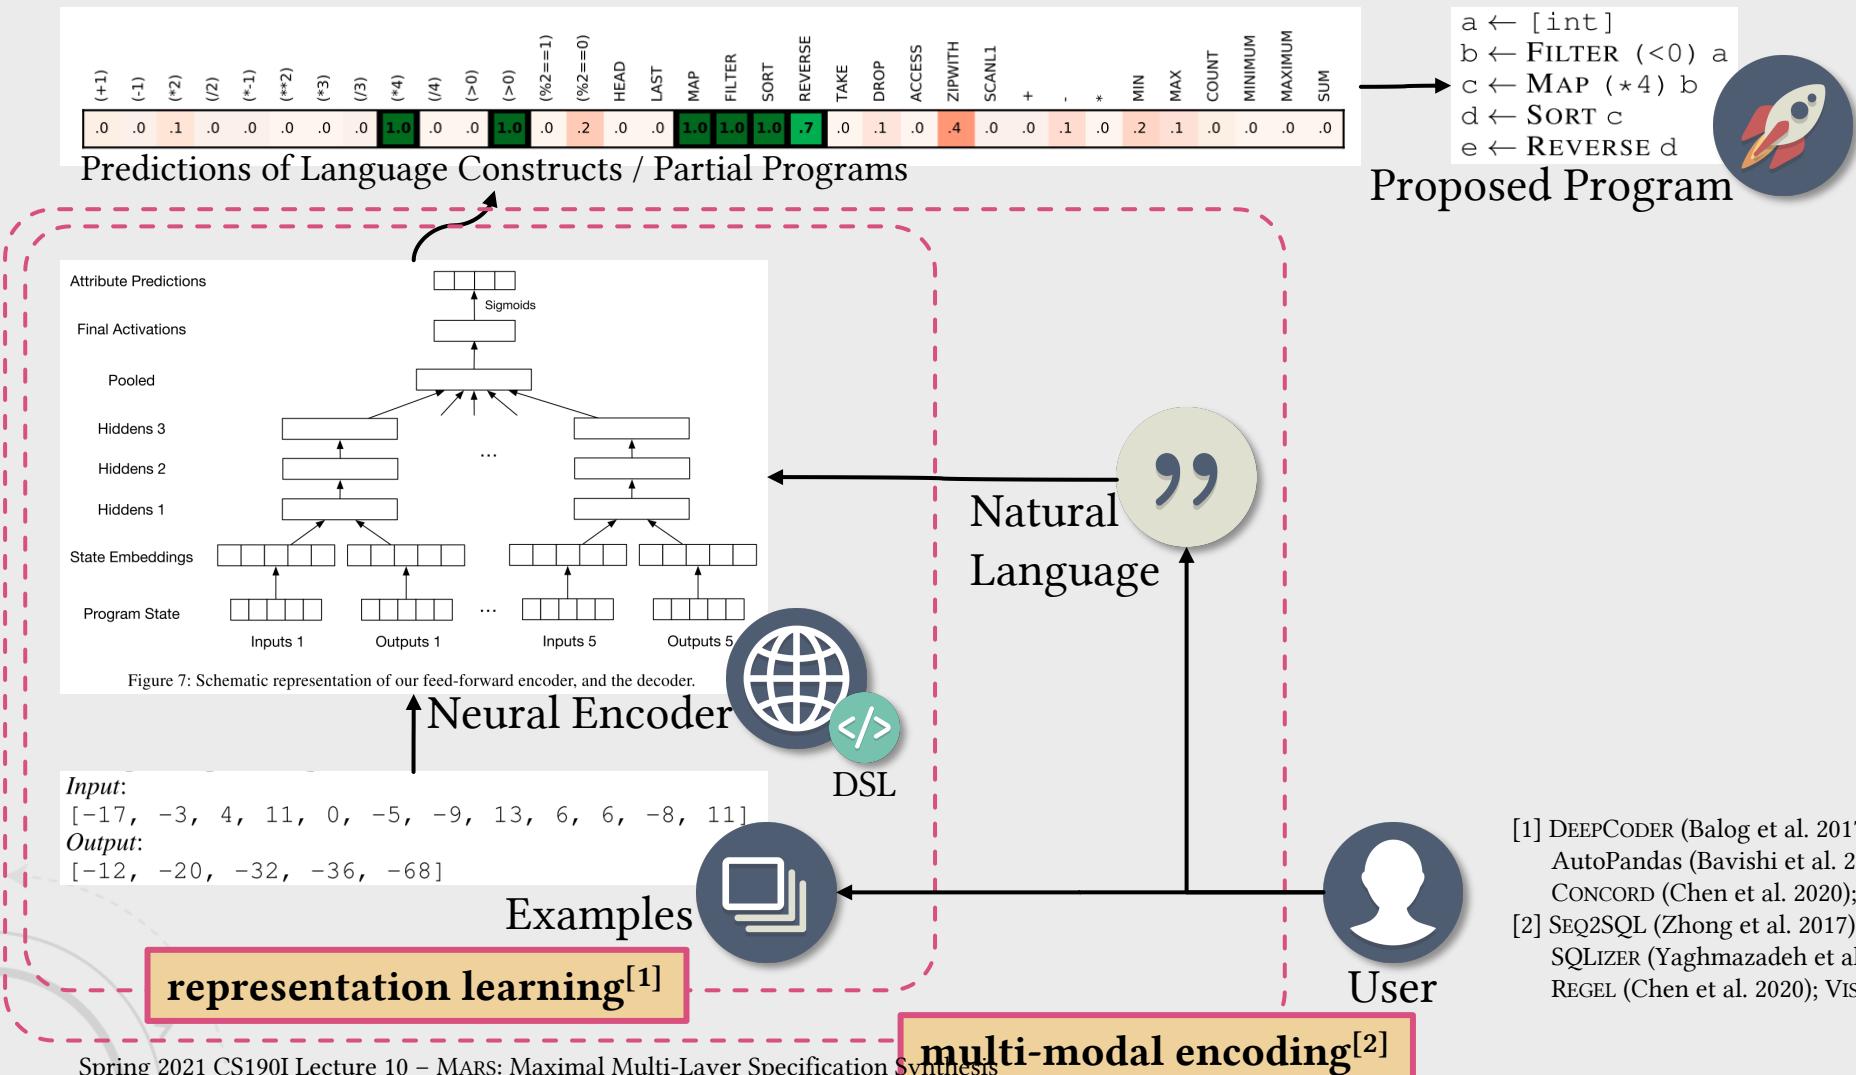
\*The table only lists some of the recent related works.

IO: Input-Output Example | NL: Natural Language | 😊: Yes | 😰: Not explicitly claimed

synthesizer	domain evaluated	specification	logical reasoning	machine learning	bridging level	multi-modal
DEEPCODER (Balog et al. 2017)	list	IO	😰	😊	NA	😰
SEQ2SQL (Zhong et al. 2017)	SQL	I + NL	😰	😊	NA	😊
DIALSQL (Gur et al. 2018)	SQL	NL	😰	😊	NA	😊
EXEC (Chen et al. 2018)	Karel	IO	😰	😊	NA	😰
NEO (Feng et al. 2018)	table + list	IO	😊	😊	★	😰
SKETCHADAPT (Nye et al. 2018)	list + string + Algolisp	IO / IO + NL	😊	😊	★★	😊
SQLIZER (Yaghmazadeh et al. 2018)	SQL	NL	😊	😊	★	😰
AutoPandas (Bavishi et al. 2019)	table	IO	😰	😊	NA	😰
<b>MARS (Chen et al. 2019)</b>	<b>table</b>	<b>IO + NL</b>	😊	😊	★★★	😊
METAL (Si et al. 2019)	circuit	logical formula	😊	😊	★★★	😰
<b>CONCORD (Chen et al. 2020)</b>	<b>list</b>	<b>IO</b>	😊	😊	★★★	😰
PROBE (Barke et al. 2020)	string + circuit + bitvector	IO	😊	😊	★★★	😰
REGEL (Chen et al. 2020)	regex	IO + NL	😊	😊	★★	😊
VISER (Wang et al. 2020)	visualization	IO + visual sketch	😊	😊	★	😊

## Program Synthesis in a Nutshell

# Program Synthesis with Machine Learning (I)



- [1] DEEPCODER (Balog et al. 2017); EXEC (Chen et al. 2018); AutoPandas (Bavishi et al. 2019); METAL (Si et al. 2019); CONCORD (Chen et al. 2020);
- [2] SEQ2SQL (Zhong et al. 2017); DIALSQL (Gur et al. 2018); SQLIZER (Yaghmazadeh et al. 2018); MARS (Chen et al. 2019); REGEL (Chen et al. 2020); VISER (Wang et al. 2020);

Program Synthesis in a Nutshell

# A Running Example from StackOverflow<sup>[1]</sup>

[Title] r script to count columns within dataset

[Example]

sample_ID	site	coll_date	species	TOT	inf_status
382870	site1	27/10/2007	SpeciesB	1	positive
382872	site2	27/10/2007	SpeciesB	1	negative
487405	site3	28/10/2007	SpeciesA	1	positive
487405	site3	28/10/2007	SpeciesA	1	positive



site	cat	sts
site1	SpeciesB_positive	1
site2	SpeciesB_negative	1
site3	SpeciesA_positive	2

[Description]

I need to reformat the data so that there is just one row per site visit (i.e. in a given site name and date combo) with columns for total found by species and the fish status (i.e. speciesA\_pos, SpeciesA\_neg, Sp\_B\_pos.. etc).

figured I need to sum within site. My thoughts were to use split/apply/aggregate/for loops etc but tried various combinations and not getting anywhere. apologies I'm not familiar with R. any comments appreciated!

[1] Example adapted from <https://stackoverflow.com/questions/39369502/r-script-to-reshape-and-count-columns-within-dataset>

## Program Synthesis in a Nutshell

# A Running DSL for Data Wrangling<sup>[1]</sup>

$t \rightarrow$	$x_i$	(input table)
	<code>select(<math>t, \vec{c}_{arg}</math>)</code>	(column projection)
	<code>unite(<math>t, c_{tgt}, \vec{c}_{arg}</math>)</code>	(column merging)
	<code>separate(<math>t, \vec{c}_{tgt}, c_{arg}</math>)</code>	(column splitting)
	<code>mutate(<math>t, c_{tgt}, op, \vec{c}_{arg}</math>)</code>	(column arithmetic)
	<code>group_by(<math>t, \vec{c}_{arg}</math>)</code>	(row grouping)
	<code>summarise(<math>t, c_{tgt}, a, \vec{c}_{arg}</math>)</code>	(row aggregation)
	<code>filter(<math>t, f, \vec{c}_{arg}</math>)</code>	(row filtering)
$op \rightarrow$	$+   -   \times   \div$	
$a \rightarrow$	$\min   \max   \text{sum}   \text{count}   \text{avg}$	

$x_i$ : the  $i$ -th input table

$t$ : table

$c, \vec{c}$ : column(s) of table

$op$ : arithmetic operation

$a$ : aggregation function

$f$ : higher-order boolean function

(input table)
(column projection)
(column merging)
(column splitting)
(column arithmetic)
(row grouping)
(row aggregation)
(row filtering)

A	B	C	D
A1	B1	1	5
A2	B2	2	6
A3	B3	3	7
A4	B4	4	8

`select`

A	C
A1	1
A2	2
A3	3
A4	4

A	B	C	D
A1	B1	1	5
A2	B2	2	6
A3	B3	3	7
A4	B4	4	8

`unite`

A_B	C	D
A1_B1	1	5
A2_B2	2	6
A3_B3	3	7
A4_B4	4	8

`separate`

A	B	C	D
A1	B1	1	5
A2	B2	2	6
A3	B3	3	7
A4	B4	4	8

`mutate`

A	B	C	D	C+D
A1	B1	1	5	6
A2	B2	2	6	8
A3	B3	3	7	10
A4	B4	4	8	12

A	B	C	D
X	B1	1	5
X	B2	2	6
Y	B3	3	7
Y	B4	4	8

`group_by`

A	avg.D
X	5.5
Y	7.5

`summarise`

A	B	C	D
A1	B1	1	5
A2	B2	2	6
A3	B3	3	7
A4	B4	4	8

A	B	C	D
A1	B1	1	5
A2	B2	2	6
A3	B3	3	7

A	B	C	D
A1	B1	1	5
A2	B2	2	6
A3	B3	3	7

[1] DSL adapted from Wang, C. et al. Visualization by Example. POPL'20

Spring 2021 CS190I Lecture 10 – MARS: Maximal Multi-Layer Specification Synthesis

## Program Synthesis in a Nutshell

# A Running Example from StackOverflow

[Example]

sample_ID	site	coll_date	species	TOT	inf_status
382870	site1	27/10/2007	SpeciesB	1	positive
382872	site2	27/10/2007	SpeciesB	1	negative
487405	site3	28/10/2007	SpeciesA	1	positive
487405	site3	28/10/2007	SpeciesA	1	positive

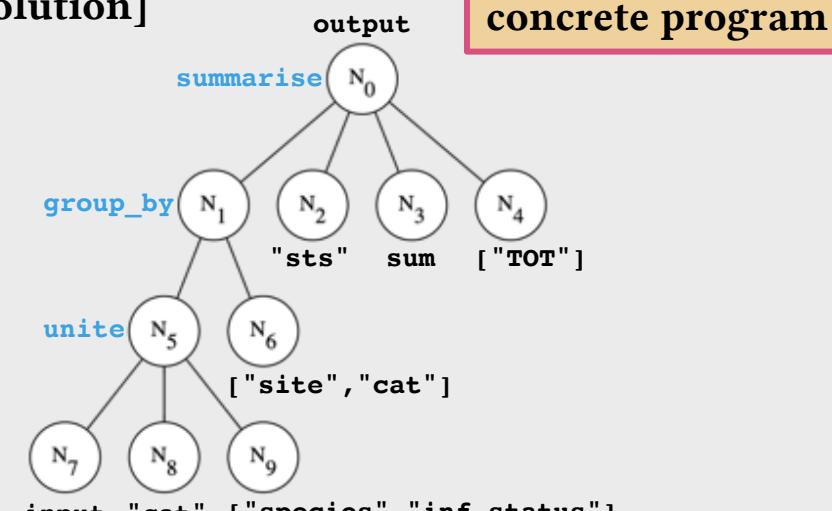
unite

sample_ID	site	coll_date	cat	TOT
382870	site1	27/10/2007	SpeciesB_positive	1
382872	site2	27/10/2007	SpeciesB_negative	1
487405	site3	28/10/2007	SpeciesA_positive	1
487405	site3	28/10/2007	SpeciesA_positive	1

group\_by  
summarise

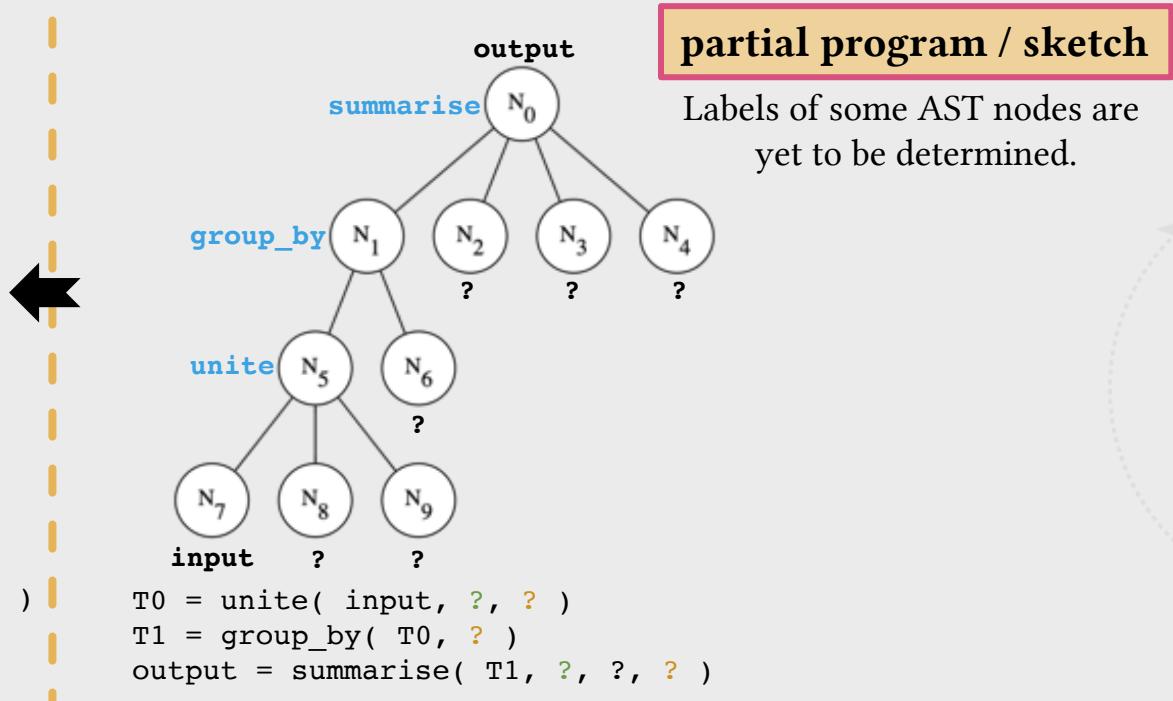
site	cat	sts
site1	SpeciesB_positive	1
site2	SpeciesB_negative	1
site3	SpeciesA_positive	2

[Solution]



```

T0 = unite( input, "cat", ["species", "inf_status"] )
T1 = group_by( T0, ["site", "cat"] )
output = summarise( T1, "sts", sum, ["TOT"] )
  
```

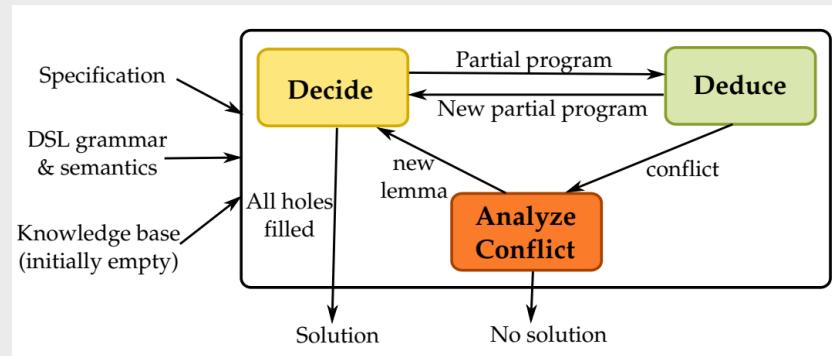


```

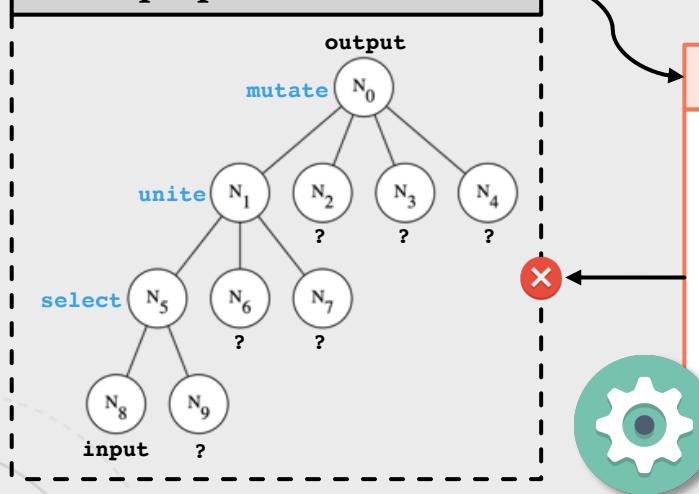
T0 = unite( input, ?, ? )
T1 = group_by( T0, ? )
output = summarise( T1, ?, ?, ? )
  
```

Program Synthesis in a Nutshell > Deductive Program Synthesis

# NEO<sup>[1]</sup>: A Brief Overview



**proposed sketch**



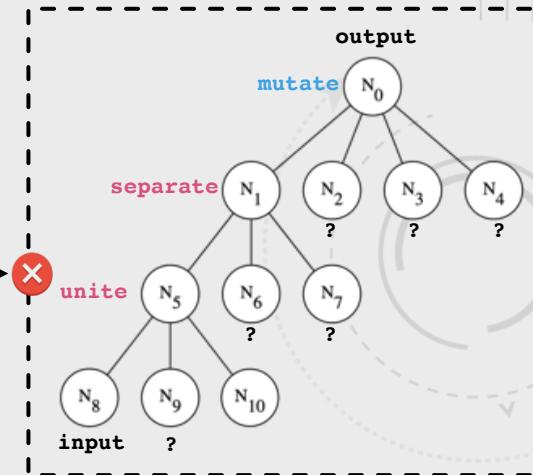
**generated constraints**

$\text{input.row} == 4 \wedge \text{input.col} == 6 \wedge$   
 $\text{N}_8.\text{row} == \text{input.row} \wedge \text{N}_8.\text{col} == \text{input.col} \wedge$   
 $\text{N}_5.\text{row} == \text{N}_8.\text{row} \wedge \text{N}_5.\text{col} <= \text{N}_8.\text{col}-1 \wedge$   
 $\text{N}_1.\text{row} == \text{N}_5.\text{row} \wedge \text{N}_1.\text{col} == \text{N}_5.\text{col}-1 \wedge$   
 $\text{N}_0.\text{row} == \text{N}_1.\text{row} \wedge \text{N}_0.\text{col} == \text{N}_1.\text{col}+1 \wedge$   
 $\text{output.row} == \text{N}_0.\text{row} \wedge \text{output.col} == \text{N}_0.\text{col} \wedge$   
 $\text{output.row} == 3 \wedge \text{output.col} == 3$

## Equivalent Modulo Conflict (EMC)<sup>[1]</sup>

<b>select</b>	$\text{out.row} == \text{in.row} \wedge \text{out.col} <= \text{in.col}-1$
<b>unite</b>	$\text{out.row} == \text{in.row} \wedge \text{out.col} == \text{in.col}-1$
<b>separate</b>	$\text{out.row} == \text{in.row} \wedge \text{out.col} == \text{in.col}+1$
<b>mutate</b>	$\text{out.row} == \text{in.row} \wedge \text{out.col} == \text{in.col}+1$
<b>group_by</b>	$\text{out.row} == \text{in.row} \wedge \text{out.col} == \text{in.col}$
<b>summarise</b>	$\text{out.row} <= \text{in.row} \wedge \text{out.col} <= \text{in.col}+1$
<b>filter</b>	$\text{out.row} <= \text{in.row}-1 \wedge \text{out.col} == \text{in.col}$

Component-Based Specifications<sup>[2]</sup> for Data Wrangling DSL



SMT-Based Deduction<sup>[1]</sup> & Analyze Conflicts<sup>[2]</sup>

[1] Feng, Y. et al. Program Synthesis using Conflict-Driven Learning. PLDI'18

[2] Feng, Y. et al. Component-based Synthesis of Table Consolidation and Transformation Tasks from Examples. PLDI'17

# Observations & Motivations

- Q1: Why logical reasoning?
  - Example: EXEC<sup>[1]</sup>
    - Concrete interpretation is less efficient, especially for complex problems
    - Logical reasoning results generalize better in pruning search space
- Q2: Why machine learning?
  - Example: AutoPandas<sup>[2]</sup>
    - Machine learning backend provides better estimations prioritizing search order
- Q3: Why bridging?
  - Example: NEO<sup>[3]</sup>
    - Programs are precise, but specifications can be vague
    - Statistical components can't reflect deduction feedbacks on the fly

We need both, and better!

[1] Chen, X. et al. Execution-Guided Neural Program Synthesis. ICLR'18

[2] Bavishi, R. et al. AutoPandas: Neural-backed Generators for Program Synthesis. OOPSLA'19

[3] Feng, Y. et al. Program Synthesis using Conflict-Driven Learning. PLDI'18

# Bridging the Logical and Statistical Lands

- Observations & Motivations
  - Existing tools do have logical and statistical components combined
    - Example: NEO<sup>[1]</sup> / TRINITY<sup>[2]</sup>
    - But they are no more than "wired" together: still talk in different languages, act independently
- Two Bridging Directions
  - MARS<sup>[3]</sup>: Encode multi-layer specifications (via machine learning) into logical components
    - Talk in logical language!
    - Encode specifications as soft/hard constraints in maximum satisfiability modulo theory (Max-SMT)
  - CONCORD<sup>[4]</sup>: Guide the statistical components using deductions
    - Talk in statistical language!
    - Generate training samples for machine learning models by explaining deduction results

[1] Feng, Y. et al. Program Synthesis using Conflict-Driven Learning. PLDI'18

[2] Martins, R. et al. Trinity: An Extensible Synthesis Framework for Data Science. VLDB'19

[3] Chen, Y. et al. Maximal Multi-layer Specification Synthesis. FSE'19

[4] Chen, Y. et al. Program Synthesis Using Deduction-Guided Reinforcement Learning. CAV'20

# MARS<sup>[1]</sup>: Encoding Multi-Layer Specifications

- Motivations
- Formalization
- Framework Overview
- Multi-Layer Specification Encoding
  - Encoding Examples as Hard Constraints
  - Encoding Natural Language Specifications
- Evaluations
  - Evaluation Setup
  - Evaluation Results & Analysis
- Discussions

[1] Chen, Y. et al. Maximal Multi-layer Specification Synthesis. FSE'19

Spring 2021 CS190I Lecture 10 – MARS: Maximal Multi-Layer Specification Synthesis

# Maximal Multi-Layer Specification Synthesis

- Motivations
  - Examples can be imprecise
  - Multi-modal specifications contain more useful information

[Title] r script to count columns within dataset

[Example]

sample_ID	site	coll_date	species	TOT	inf_status
382870	site1	27/10/2007	SpeciesB	1	positive
382872	site2	27/10/2007	SpeciesB	1	negative
487405	site3	28/10/2007	SpeciesA	1	positive
487405	site3	28/10/2007	SpeciesA	1	positive



site	cat	sts
site1	SpeciesB_positive	1
site2	SpeciesB_negative	1
site3	SpeciesA_positive	2

[Description]

I need to reformat the data so that there is just one row per site visit (i.e. in a given site name and date combo) with columns for total found by species and the fish status (i.e. speciesA\_pos, SpeciesA\_neg, Sp\_B\_pos.. etc).

figured I need to sum within site. My thoughts were to use split/apply/aggregate/for loops etc but tried various combinations and not getting anywhere. apologies I'm not familiar with R. any comments appreciated!

# Maximal Multi-Layer Specification Synthesis

- Motivations
  - Examples can be imprecise
  - Multi-modal specifications contain more useful information

[Title] r script to count columns within dataset

[Example]

sample_ID	site	coll_date	species	TOT	inf_status
382870	site1	27/10/2007	SpeciesB	1	positive
382872	site2	27/10/2007	SpeciesB	1	negative
487405	site3	28/10/2007	SpeciesA	1	positive
487405	site3	28/10/2007	SpeciesA	1	positive



site	cat	sts
site1	SpeciesB_positive	1
site2	SpeciesB_negative	1
site3	SpeciesA_positive	2

[Description]

I need to reformat the data so that there is just one row per site visit (i.e. in a given site name and date combo) with columns for total found by species and the fish status (i.e. speciesA\_pos, SpeciesA\_neg, Sp\_B\_pos.. etc).

figured I need to sum within site. My thoughts were to use split/apply/aggregate/for loops etc but tried various combinations and not getting anywhere. apologies I'm not familiar with R. any comments appreciated!

# Maximal Multi-Layer Specification Synthesis

- Motivations
  - Examples can be imprecise
  - Multi-modal specifications contain more useful information

[Title] r script to count columns within dataset

[Example]

sample_ID	site	coll_date	species	TOT	inf_status
382870	site1	27/10/2007	SpeciesB	1	positive
382872	site2	27/10/2007	SpeciesB	1	negative
487405	site3	28/10/2007	SpeciesA	1	positive
487405	site3	28/10/2007	SpeciesA	1	positive



site	cat	sts
site1	SpeciesB_positive	1
site2	SpeciesB_negative	1
site3	SpeciesA_positive	2

[Description]

I need to reformat the data so that there is just one row per site visit (i.e. in a given site name and date combo) with columns for total found by species and the fish status (i.e. speciesA\_pos, SpeciesA\_neg, Sp\_B\_pos.. etc).

figured I need to sum within site. My thoughts were to use split/apply/aggregate/for loops etc but tried various combinations and not getting anywhere. apologies I'm not familiar with R. any comments appreciated!

# Maximal Multi-Layer Specification Synthesis

- Motivations

- Example: `group_by` imprecise
- Multi-modal specifications contain more useful information

[Title] r script to count columns within dataset

[Example]

sample_ID	site	coll_date	species	TOT	inf_status
382870	site1	27/10/2007	SpeciesB	1	positive
382872	site2	27/10/2007	SpeciesB	1	negative
487405	site3	28/10/2007	SpeciesA	1	positive
487405	site3	28/10/2007	SpeciesA	1	positive

`unite`

site	cat	sts
site1	SpeciesB_positive	1
site2	SpeciesB_negative	1
site3	SpeciesA_positive	2

`summarise`

[Description]

I need to reformat the data so that there is just one row per site visit (i.e. in a given site name and date combo) with columns for total found by species and the fish status (i.e. `SpeciesA_pos`, `SpeciesA_neg`, `Sp_B_pos..` etc).

figured I need to sum within site. My thoughts were to use `split/apply/aggregate/for loops` etc but tried various combinations and not getting anywhere. apologies I'm not familiar with R. any comments appreciated!

# Formalization

- Maximal Multi-Layer Specification Synthesis

hard constraints/specifications: examples

Given specification  $(\mathcal{E}, \Psi, \Sigma)$  where  $\mathcal{E} = (T_{in}, T_{out})$ ,  $\Psi = \bigcup (\chi_i, \omega_i)$ , and  $\Sigma$  represents all symbols in the DSL, the *Maximal Multi-Layer Specification Synthesis* problem is to infer a program  $\mathcal{P}$  such that:

- $\mathcal{P}$  is a well-typed expression over symbols in  $\Sigma$ ,
- $\mathcal{P}(T_{in}) = T_{out}$ , and
- $\sum \omega_i$  is maximized

soft constraints/specifications: natural languages

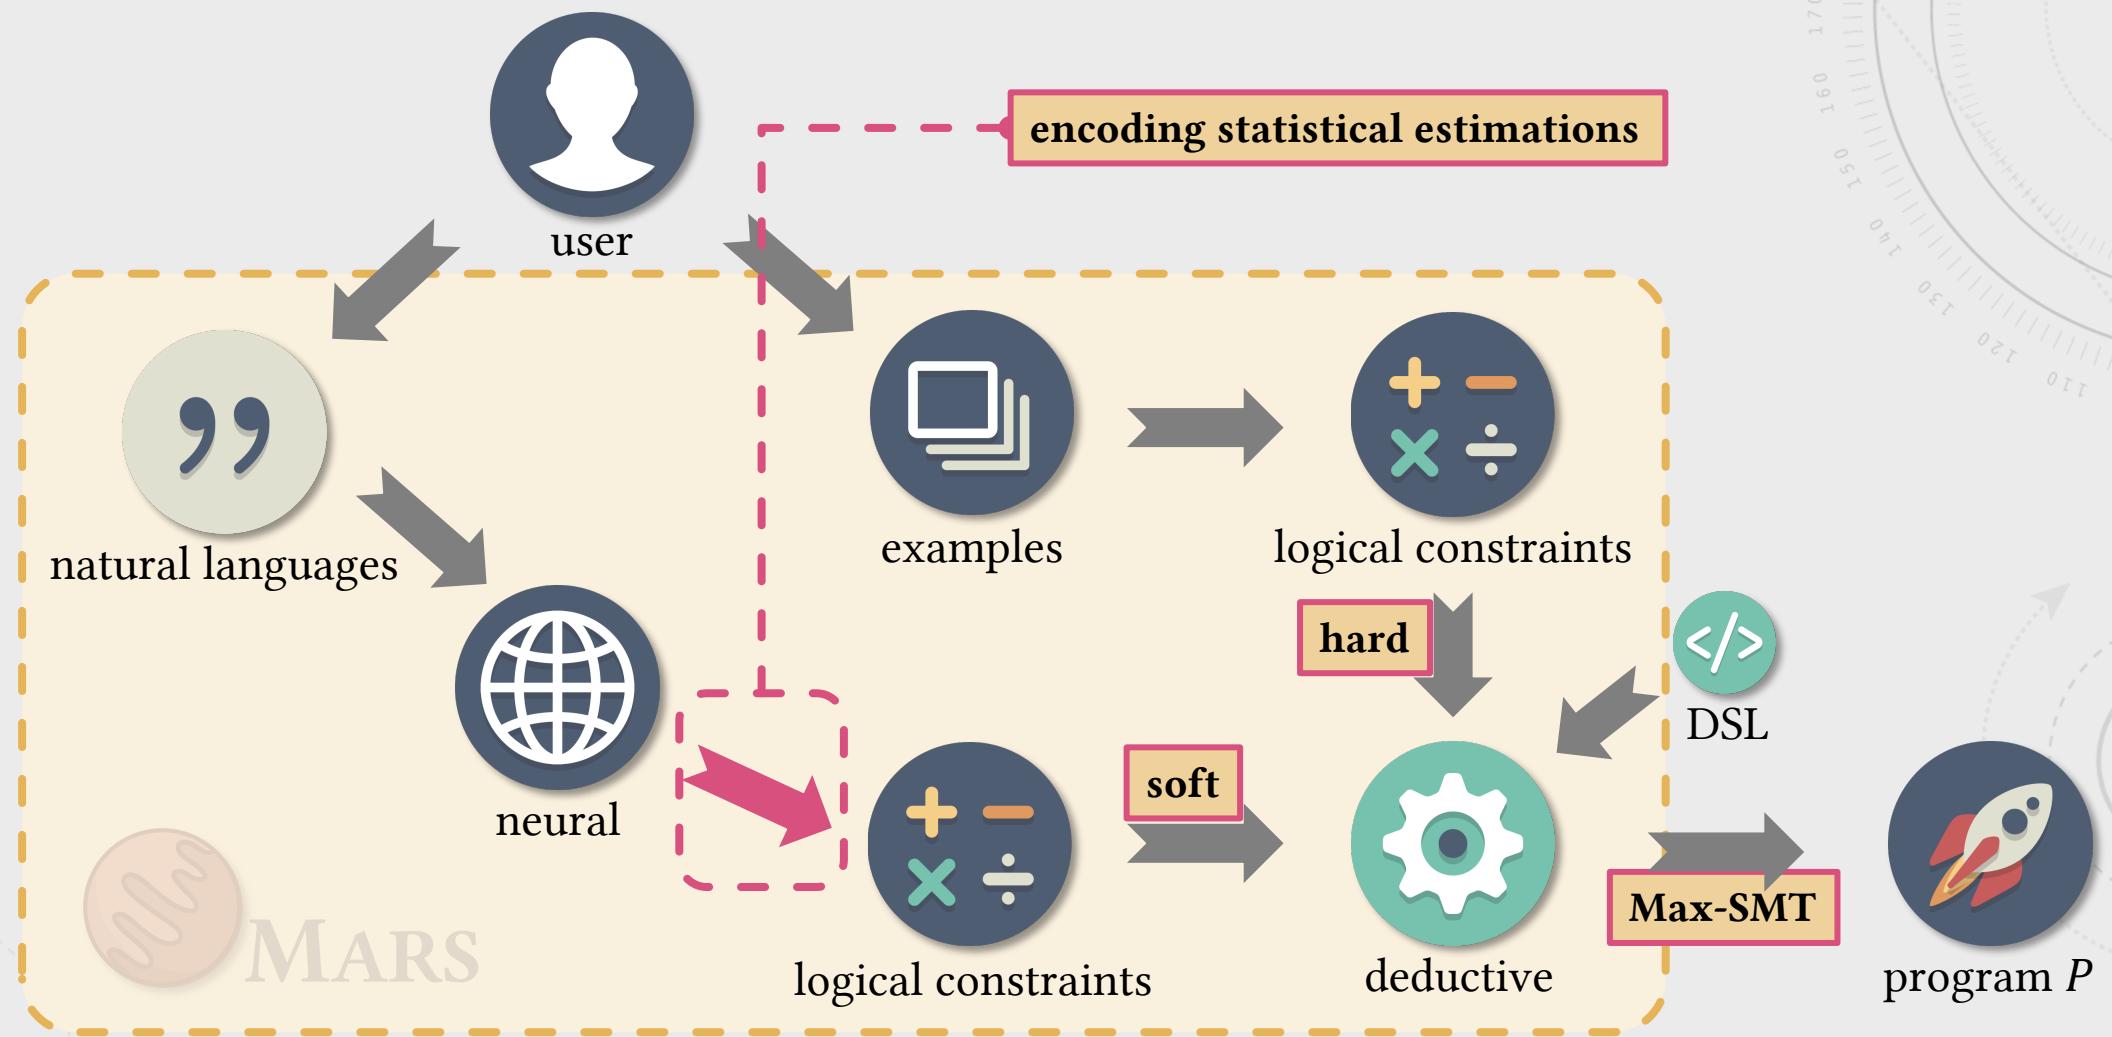
DSL construct

preference/confidence

- We model the problem using *maximum satisfiability modulo theory* (Max-SMT) and solve it with an off-the-shelf SMT solver.

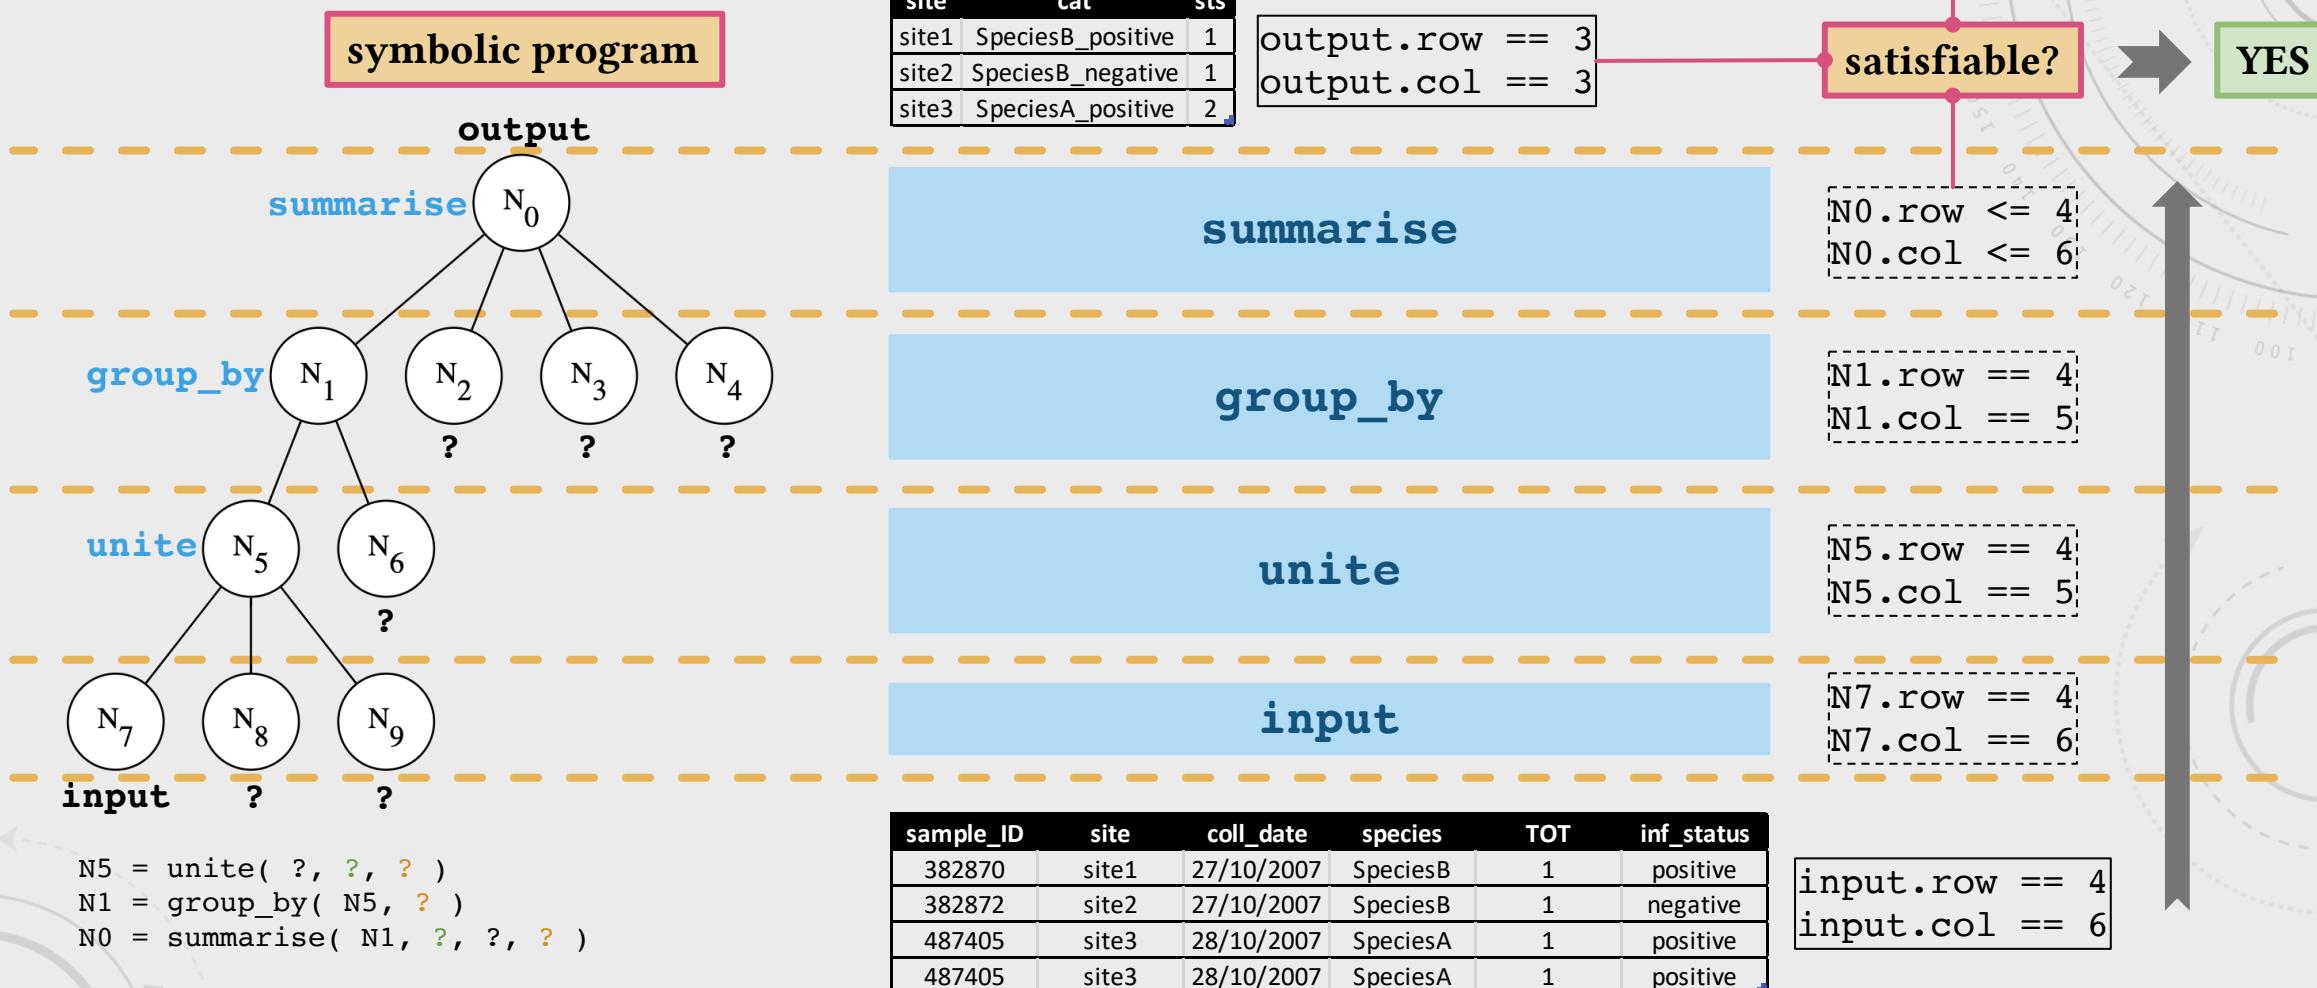
**Hard constraints should be satisfied;  
Soft constraints should be maximized.**

# Framework Overview

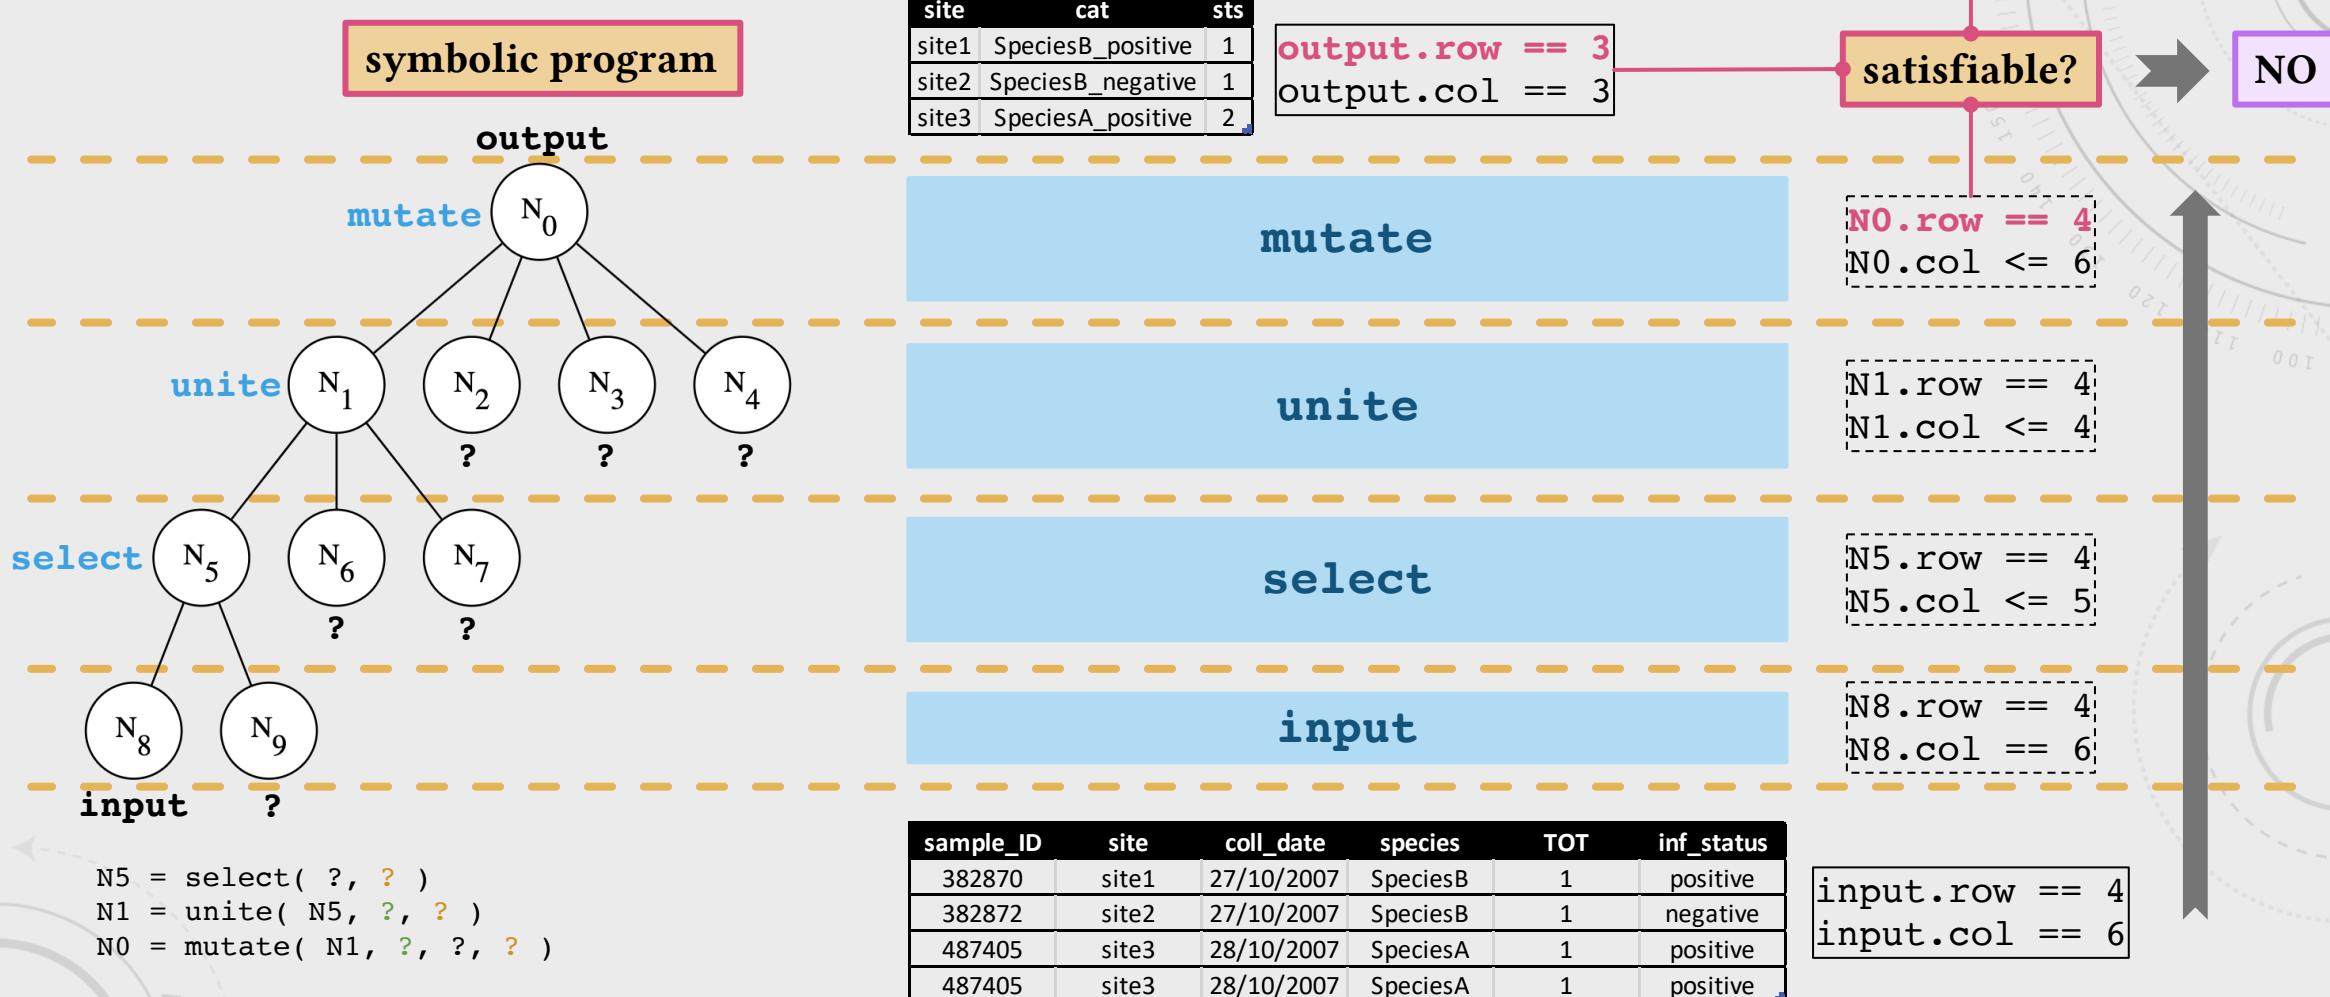


# Encoding Examples as Hard Constraints

output.row == N0.row  
output.col == N0.col

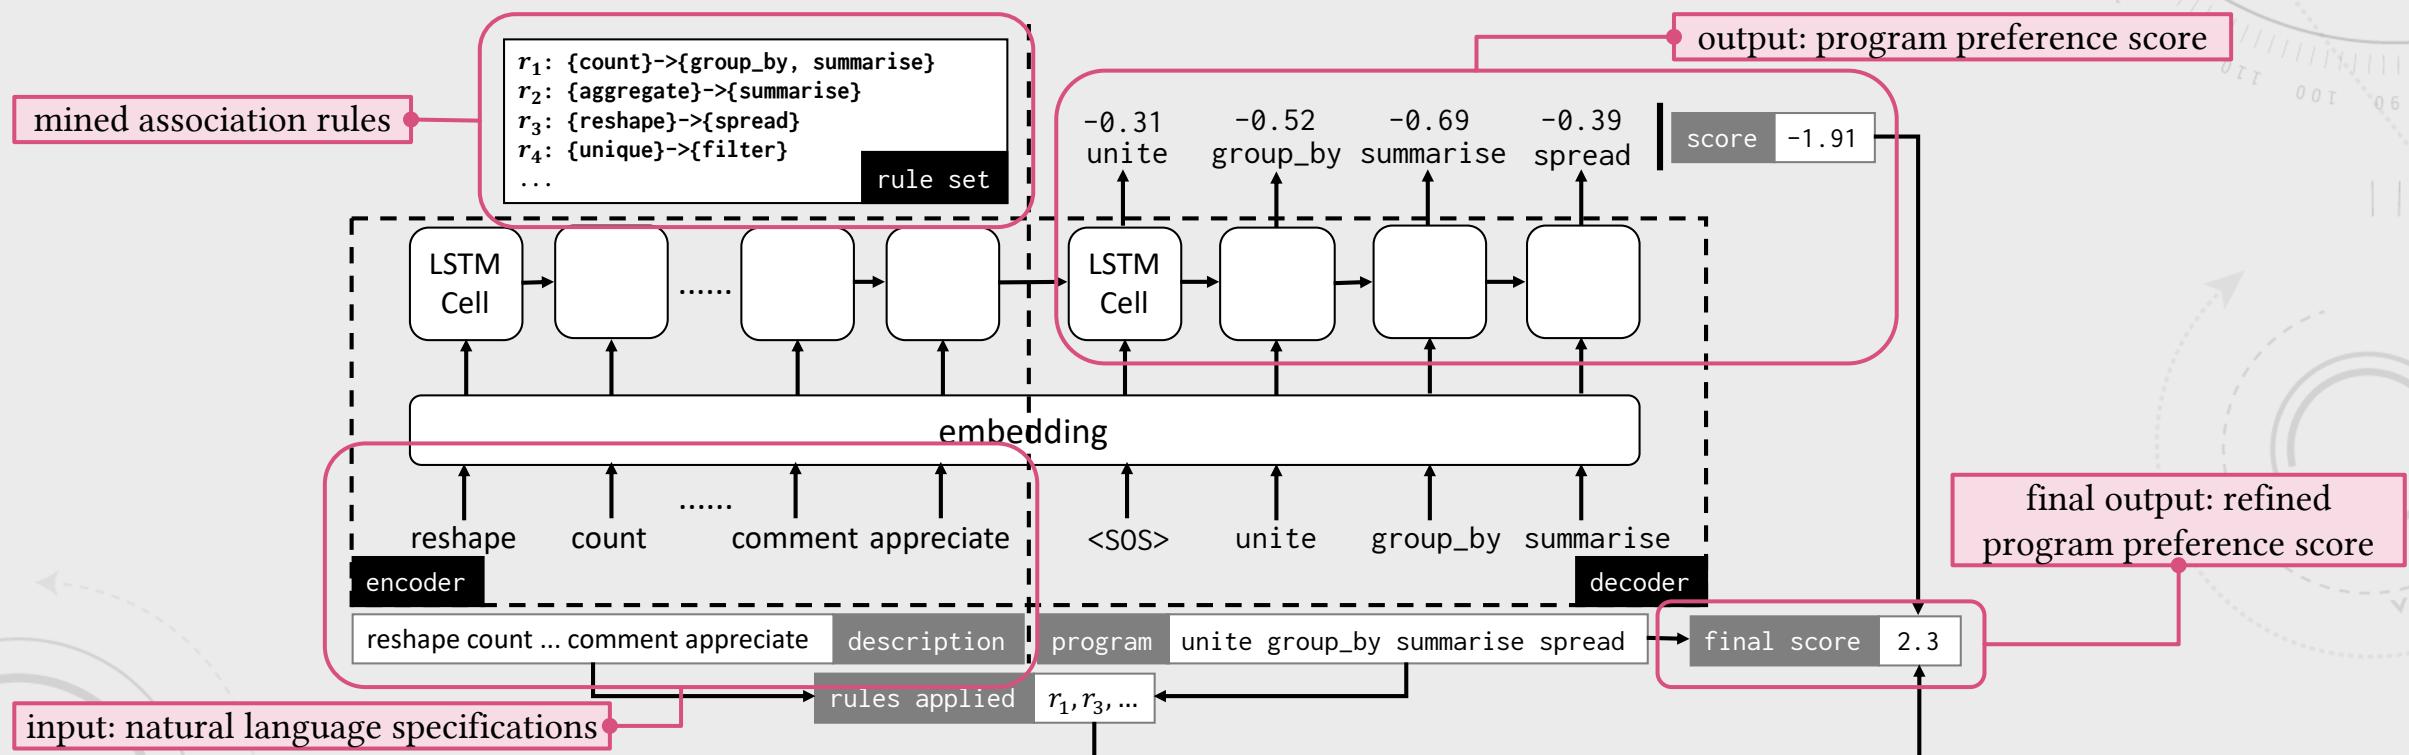


# Encoding Examples as Hard Constraints



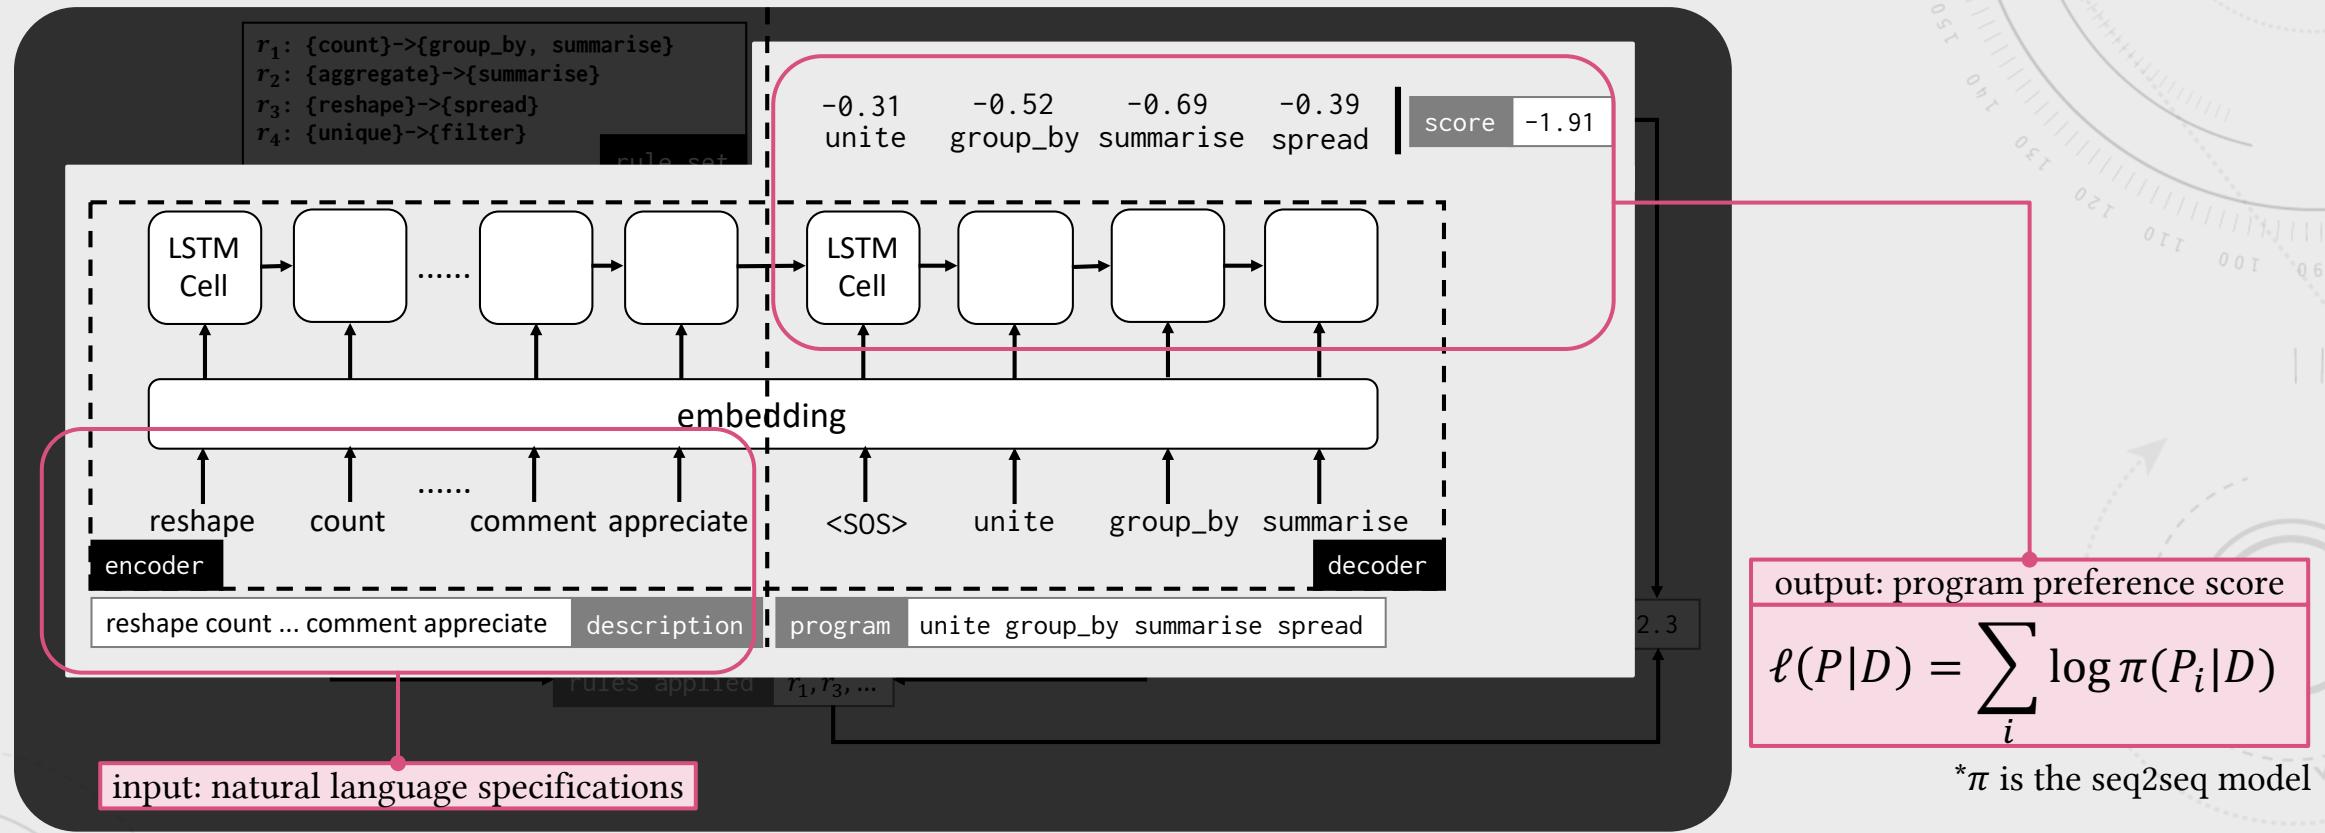
# Encoding Natural Language Specifications

- The Hybrid Neural Architecture
  - seq2seq model (supervised): capture common natural language semantics
  - association rule module (unsupervised): capture frequent patterns and refine the preference



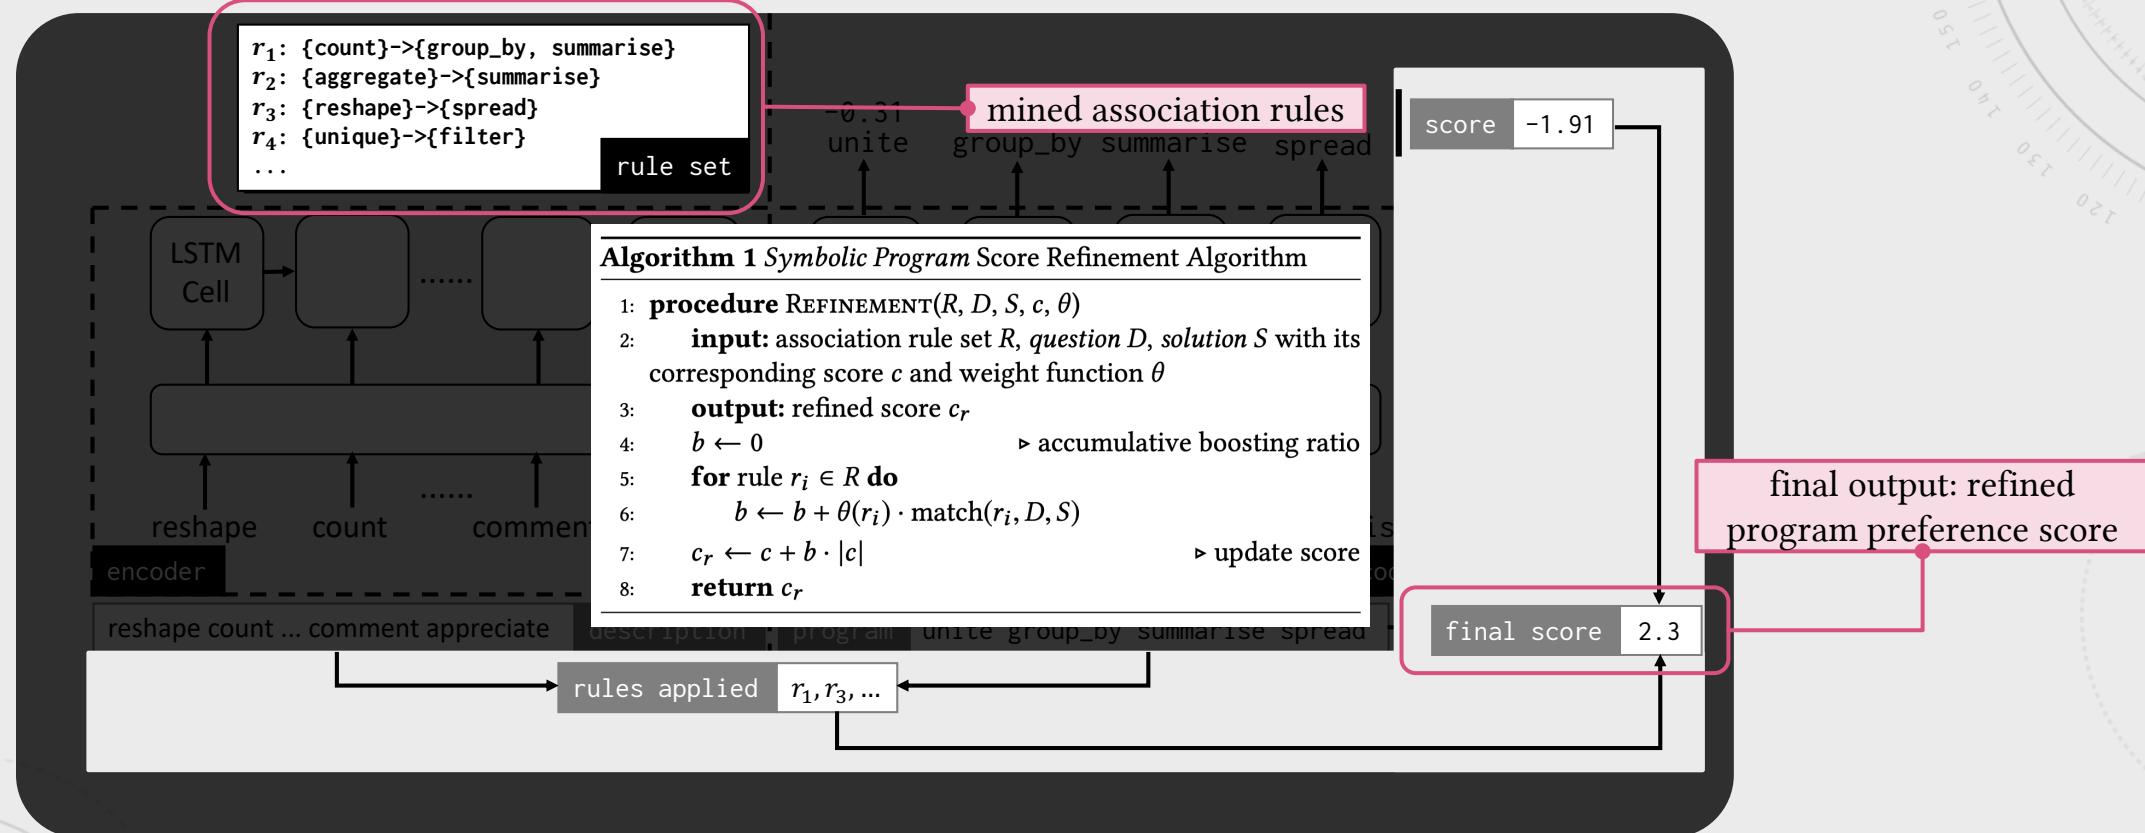
# Encoding Natural Language Specifications

- The seq2seq model



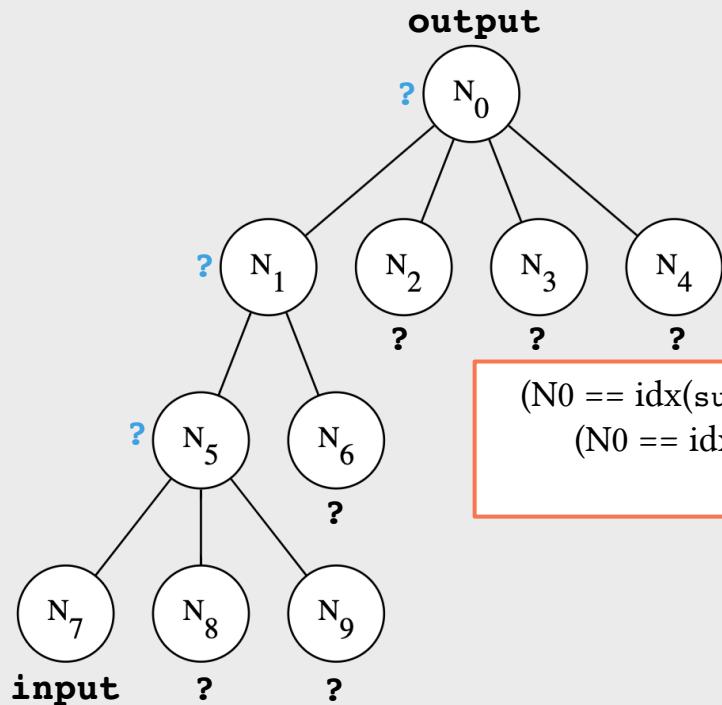
# Encoding Natural Language Specifications

- The association rule module

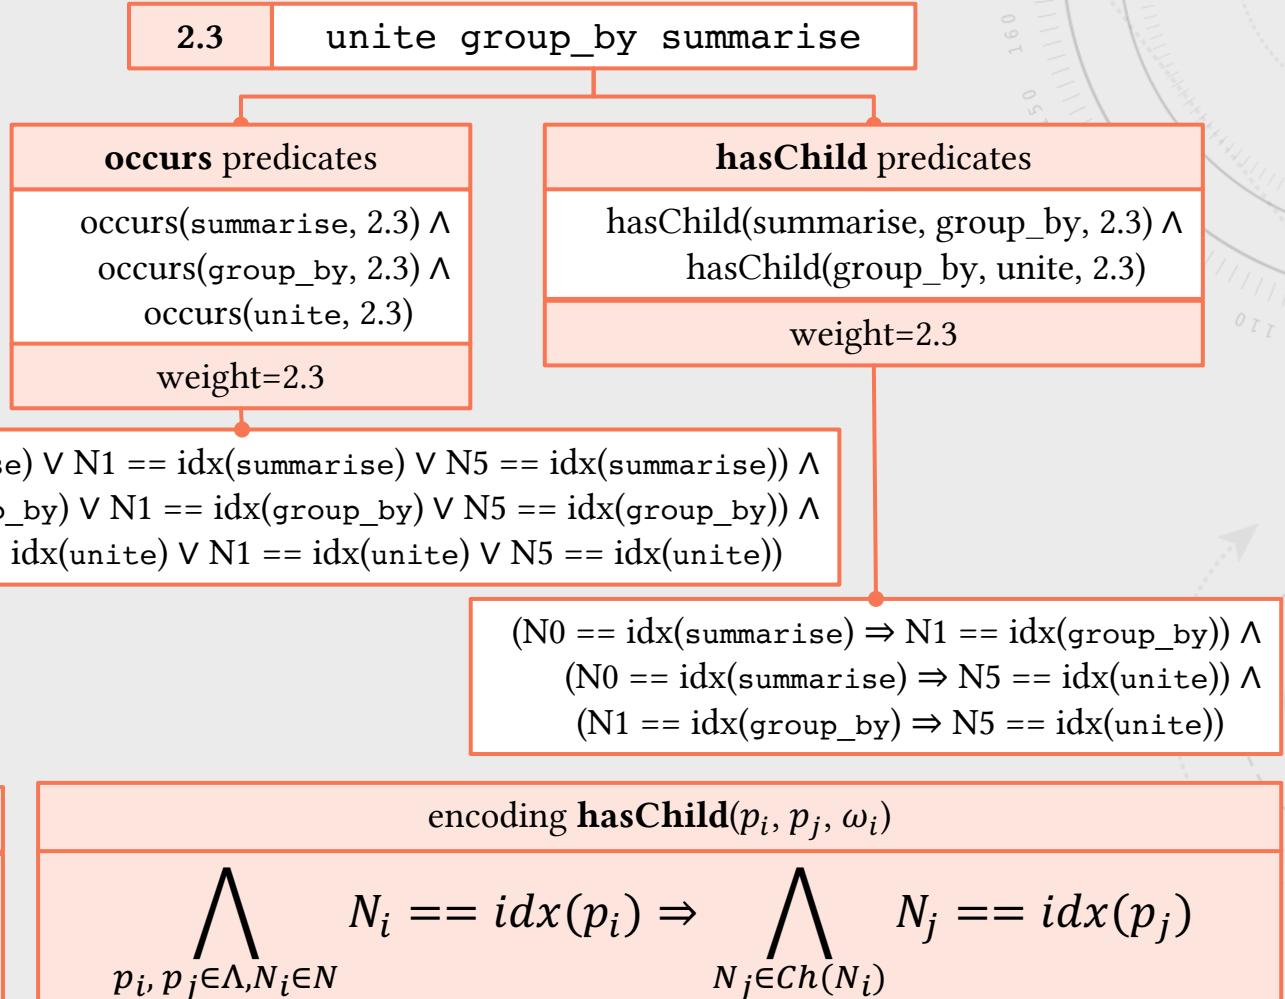


# Encoding Natural Language Specifications

- Encoding refined preference scores



encoding **occurs**( $p_i, \omega_i$ )

$$\bigwedge_{p_i \in \Lambda} \bigvee_{N_i \in N} N_i == idx(p_i)$$


# Evaluation Setup

- Research Questions
  - Q1: Do our multi-layer specification and neural architecture suggest candidates that are close to the user intent?
  - Q2: What is the impact of the neural architecture in MARS on the performance of a state-of-the-art synthesizer for data wrangling tasks?
  - Q3: How is the performance of MARS affected by the quality of the corpus?
- Experiment Setup
  - Benchmarks: 80 Real-World Challenging Data Wrangling Tasks
  - Dataset: 20,640 StackOverflow Pages of Data Wrangling Tasks
    - 16,459 question-solution pairs for seq2seq model
    - 37,748 transactions for association rule mining (*Apriori* algorithm); we obtain 187 *valid*<sup>[1]</sup> rules
  - Comparison to MORPHEUS<sup>[2]</sup>

[1] A rule is *valid* if its confidence  $\geq 0.9$  or support  $\geq 0.003$ , and satisfies all the criteria defined in Chen, Y. et al.. Maximal Multi-layer Specification Synthesis. FSE'19

[2] Feng, Y. et al.. Component-based Synthesis of Table Consolidation and Transformation Tasks from Examples. PLDI'17

# Evaluation Results & Analysis

- Timeout: 5 mins
- Ablation Variants
  - *ngram*: built-in statistical model in MOPHEUS
  - *seq2seq*: MARS with seq2seq model
  - *hybrid*: MARS with seq2seq model and preference score refinement (association rules)

**Table 1: Statistics for different model rankings.**

model	<i>n-gram</i>	<i>seq2seq</i>	<i>hybrid</i>
<b>average</b> <sup>*</sup>	42	25	18
<b>std.<sup>1</sup></b>	70	39	26

<sup>1</sup> standard deviation.

\* computed based on the rankings of the correct solutions.

**Table 2: Counts of top-1s and top-3s in different models.**

model	<i>n-gram</i>	<i>seq2seq</i>	<i>hybrid</i>
<b>Top-1 total</b> <sup>*</sup>	0	8	11
<b>Top-3 total</b> <sup>*</sup>	2	18	29

\* computed based on the rankings of the correct solutions.

**Table 3: Statistics of running time.**

model	avg. speedup <sup>1</sup>	#timeouts <sup>*</sup>
<i>ngram</i>	1x	11
<i>seq2seq</i>	6x	8
<i>hybrid</i>	15x	2

<sup>1</sup> average speedup on challenging solved benchmarks.

\* number of timeouts on all benchmarks.

# Discussions

- Limitations
  - Insufficient Text
    - Description of the question is barely useful
  - Contextual Text
    - Some questions require understanding of pragmatic contexts, not only semantic
  - Misleading Text
    - User specifies functionality not supported by the DSL
- Threats to Validity
  - Quality of the Corpus
  - Benchmark Selection

*“... I can solve my problem using dplyr’s **mutate** but it’s a time-intensive, roundabout way to achieve my goal. ...”*

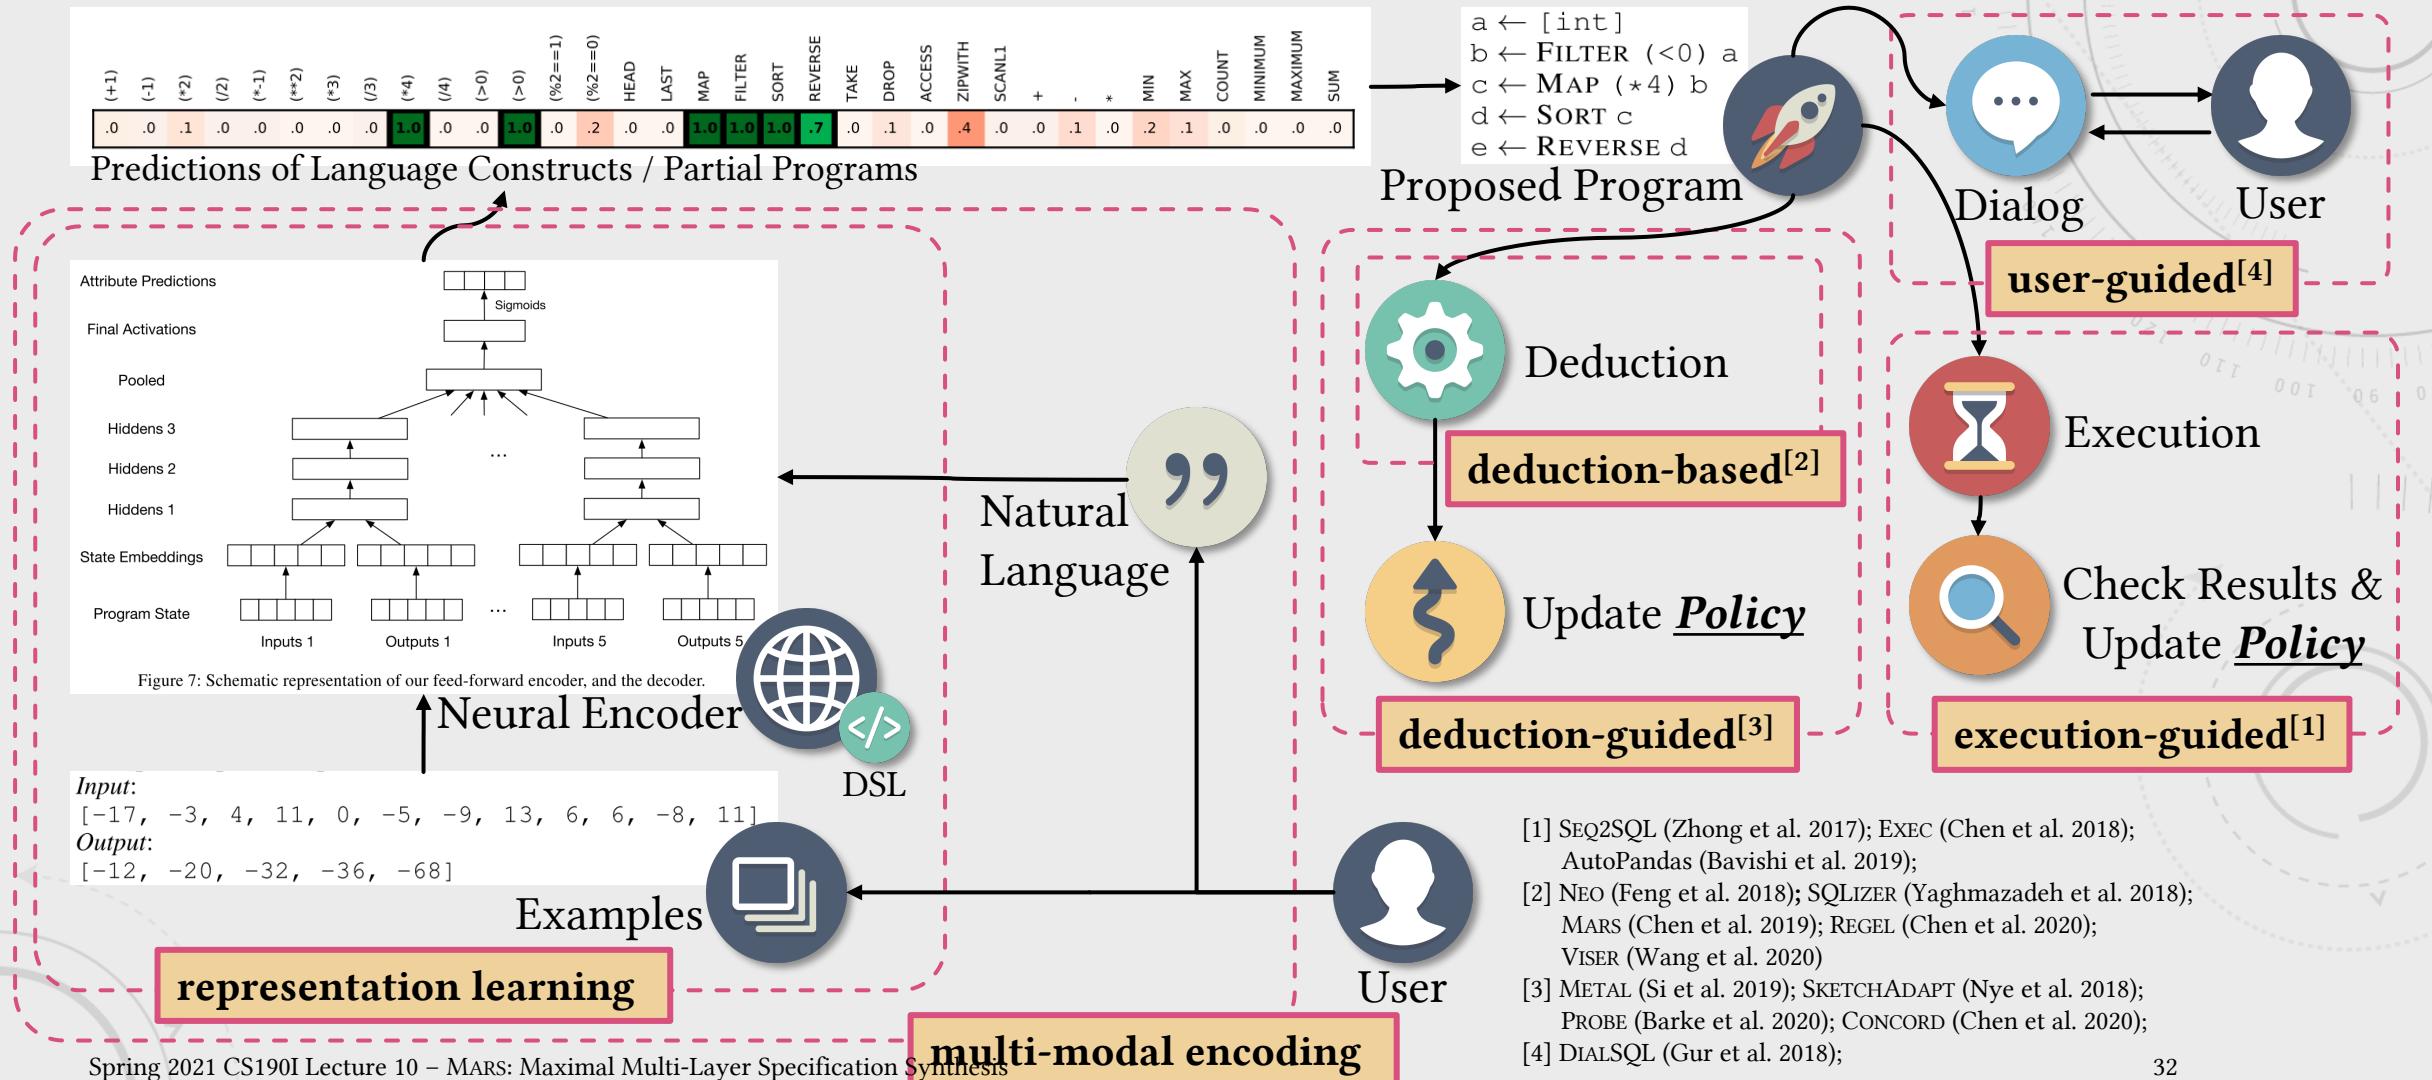
*“... I want to use **mutate** to make variable *d* which is mean of *a,b* and *c*. ...”*

# Related Works & Conclusions

- Program Synthesis with Machine Learning (II)
- Challenges, Conclusions & Future Works

## Related Works & Conclusions

# Program Synthesis with Machine Learning (II)



[1] SEQ2SQL (Zhong et al. 2017); EXEC (Chen et al. 2018); AutoPandas (Bavishi et al. 2019);

[2] NEO (Feng et al. 2018); SQLIZER (Yaghmazadeh et al. 2018); MARS (Chen et al. 2019); REGEL (Chen et al. 2020); VISER (Wang et al. 2020)

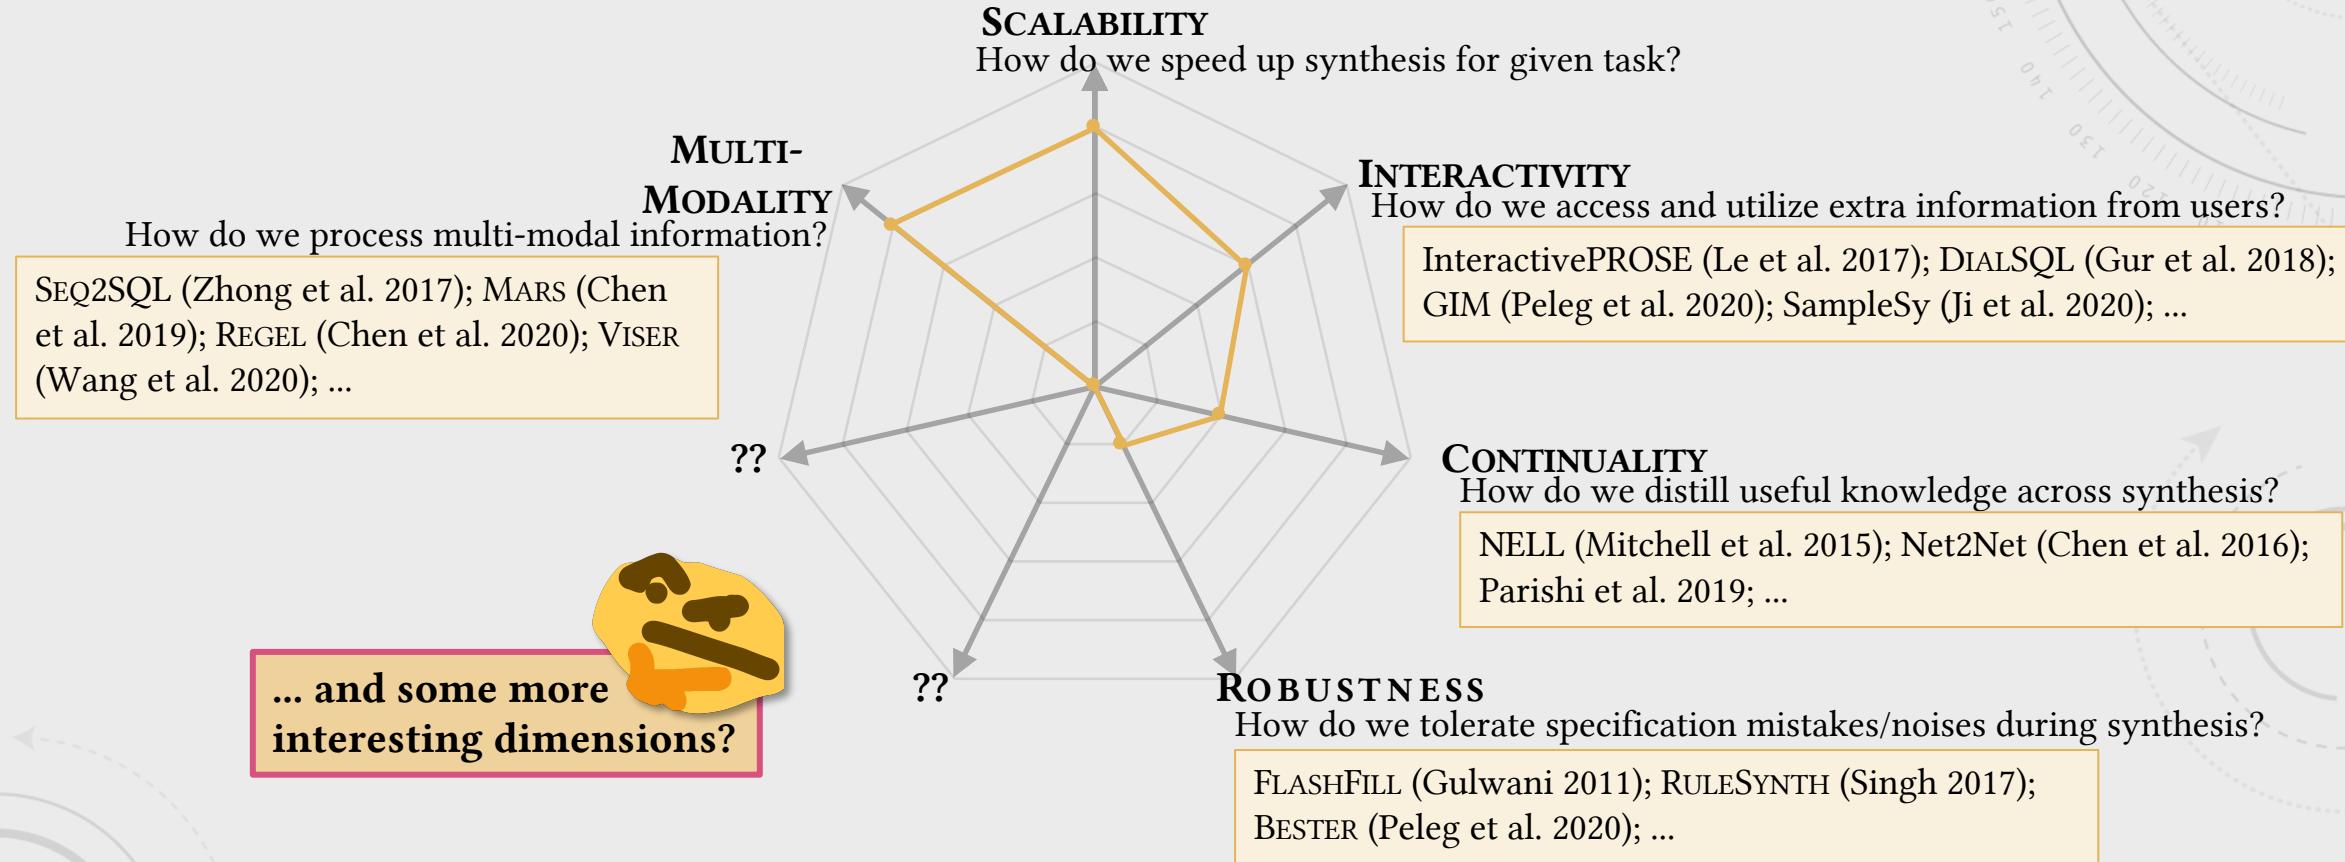
[3] METAL (Si et al. 2019); SKETCHADAPT (Nye et al. 2018); PROBE (Barke et al. 2020); CONCORD (Chen et al. 2020);

[4] DIALSQL (Gur et al. 2018);

## Related Works & Conclusions

# Challenges, Conclusions & Future Works

DEEPCODER (Balog et al. 2017); EXEC (Chen et al. 2018); NEO (Feng et al. 2018); SQLIZER (Yaghmazadeh et al. 2018); AutoPandas (Bavishi et al. 2019); METAL (Si et al. 2019); SKETCHADAPT (Nye et al. 2018); PROBE (Barke et al. 2020); CONCORD (Chen et al. 2020); ...



# References I

- Gulwani, S. Automating String Processing in Spreadsheets using Input-Output Examples. In POPL'11
- Berant, J., Chou, A., Frostig, R., & Liang, P. Semantic Parsing on {F}reebase from Question-Answer Pairs. In EMNLP'13
- Mitchell, T., Cohen, W., Hruschka, E., Talukdar, P., Betteridge, J., Carlson, A., ... Welling, J. Never-Ending Learning. In AAAI'15
- Chen, T., Goodfellow, I. J., & Shlens, J. Net2Net: Accelerating Learning via Knowledge Transfer. In ICLR'16
- Balog, M., Gaunt, A. L., Brockschmidt, M., Nowozin, S., & Tarlow, D. DeepCoder: Learning to Write Programs. In ICLR'17
- Feng, Y., Martins, R., Van Geffen, J., Dillig, I., & Chaudhuri, S. Component-based Synthesis of Table Consolidation and Transformation Tasks from Examples. In PLDI'17
- Le, V., Perelman, D., Polozov, O., Raza, M., Udupa, A., & Gulwani, S. Interactive Program Synthesis. CoRR, abs/1703.0
- Singh, R., Meduri, V. V., Elmagarmid, A., Madden, S., Papotti, P., Quiané-Ruiz, J.-A., ... Tang, N. Synthesizing Entity Matching Rules by Examples. In VLDB'17
- Yaghmazadeh, N., Wang, Y., Dillig, I., & Dillig, T. SQLizer: Query Synthesis from Natural Language. In OOPSLA'17
- Zhong, V., Xiong, C., & Socher, R. Seq2SQL: Generating Structured Queries from Natural Language using Reinforcement Learning. CoRR, abs/1709.0
- Dong, L., & Lapata, M. Coarse-to-Fine Decoding for Neural Semantic Parsing. In ACL'18
- Feng, Y., Martins, R., Bastani, O., & Dillig, I. Program Synthesis using Conflict-Driven Learning. In PLDI'18
- Gur, I., Yavuz, S., Su, Y., & Yan, X. {D}ial{SQL}: Dialogue Based Structured Query Generation. In ACL'18
- Peleg, H., Shoham, S., & Yahav, E. Programming Not Only by Example. In ICSE'18
- Bavishi, R., Lemieux, C., Fox, R., Sen, K., & Stoica, I. AutoPandas: Neural-backed Generators for Program Synthesis. In OOPSLA'19

# References II

- Chen, X., Liu, C., & Song, D. Execution-Guided Neural Program Synthesis. In ICLR'19
- **Chen, Y., Martins, R., & Feng, Y. Maximal Multi-layer Specification Synthesis. In FSE'19**
- Dai, W.-Z., Xu, Q., Yu, Y., & Zhou, Z.-H. Bridging Machine Learning and Logical Reasoning by Abductive Learning. In NeurIPS'19
- **Martins, R., Chen, J., Chen, Y., Feng, Y., & Dillig, I. Trinity: An Extensible Synthesis Framework for Data Science. In VLDB'19**
- Nye, M., Hewitt, L., Tenenbaum, J., & Solar-Lezama, A. Learning to Infer Program Sketches. In ICML'19
- Parisi, G. I., Kemker, R., Part, J. L., Kanan, C., & Wermter, S. Continual lifelong learning with neural networks: A review. *Neural Networks*, 113, 54–71.
- Si, X., Yang, Y., Dai, H., Naik, M., & Song, L. Learning a Meta-Solver for Syntax-Guided Program Synthesis. In ICLR'19
- Barke, S., Peleg, H., & Polikarpova, N. Just-in-Time Learning for Bottom-up Enumerative Synthesis. In OOPSLA'20
- Chen, Q., Wang, X., Ye, X., Durrett, G., & Dillig, I. Multi-Modal Synthesis of Regular Expressions. In PLDI'20
- **Chen, Y., Wang, C., Bastani, O., Dillig, I., & Feng, Y. Program Synthesis Using Deduction-Guided Reinforcement Learning. In CAV'20**
- Ji, R., Liang, J., Xiong, Y., Zhang, L., & Hu, Z. Question Selection for Interactive Program Synthesis. In PLDI'20
- **Mariano, B., Chen, Y., Feng, Y., Lahiri, S., & Dillig, I. Demystifying Loops in Smart Contracts. In ASE'20**
- Peleg, H., & Polikarpova, N. Perfect is the Enemy of Good: Best-Effort Program Synthesis. In ECOOP'20
- Wang, C., Feng, Y., Bodik, R., Cheung, A., & Dillig, I. Visualization by Example. In POPL'20