

Lecture 11: Synthesis with Stochastic Search

Yu Feng
Spring 2021

Summary of previous lecture

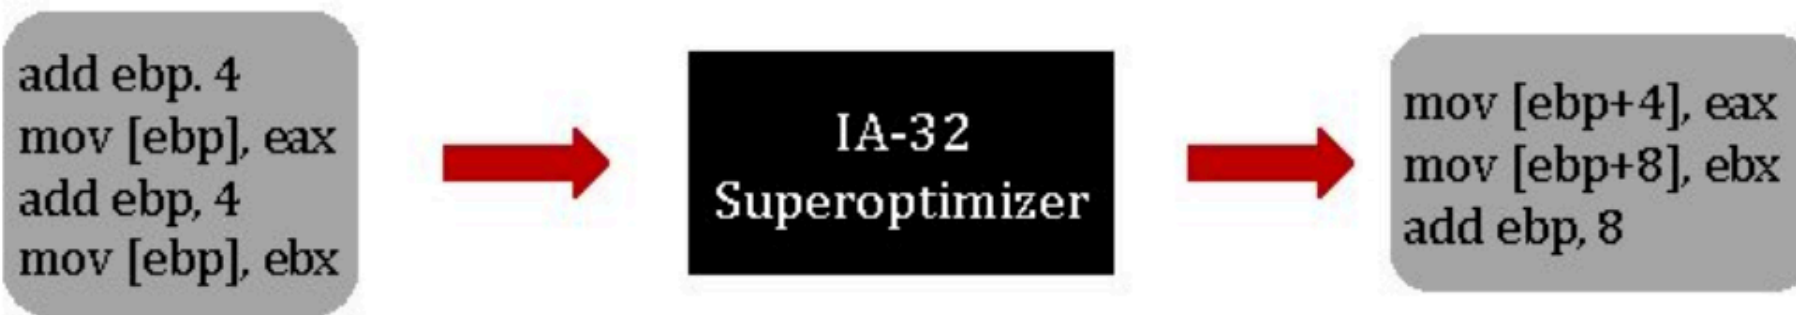
- 2nd homework is due today
- Synthesis with machine learning

Performance

- Many systems where code performance matters
 - Compute-bound
 - Repeatedly executed
- Scientific computing
- Graphics
- Low-latency server code
- Encryption/decryption



Superoptimization



Superoptimization

- Generate the optimal implementation
 - Massalin [ASPLOS'87]
 - Enumerate all possible straight-line programs
- Stoke, Schkufza [ASPLOS'13]
 - Random enumeration instead of exhaustive
 - Transforms programs with loops

$$[r8:rdi] = rsi * [ecx:edx] + r8 + rdi$$

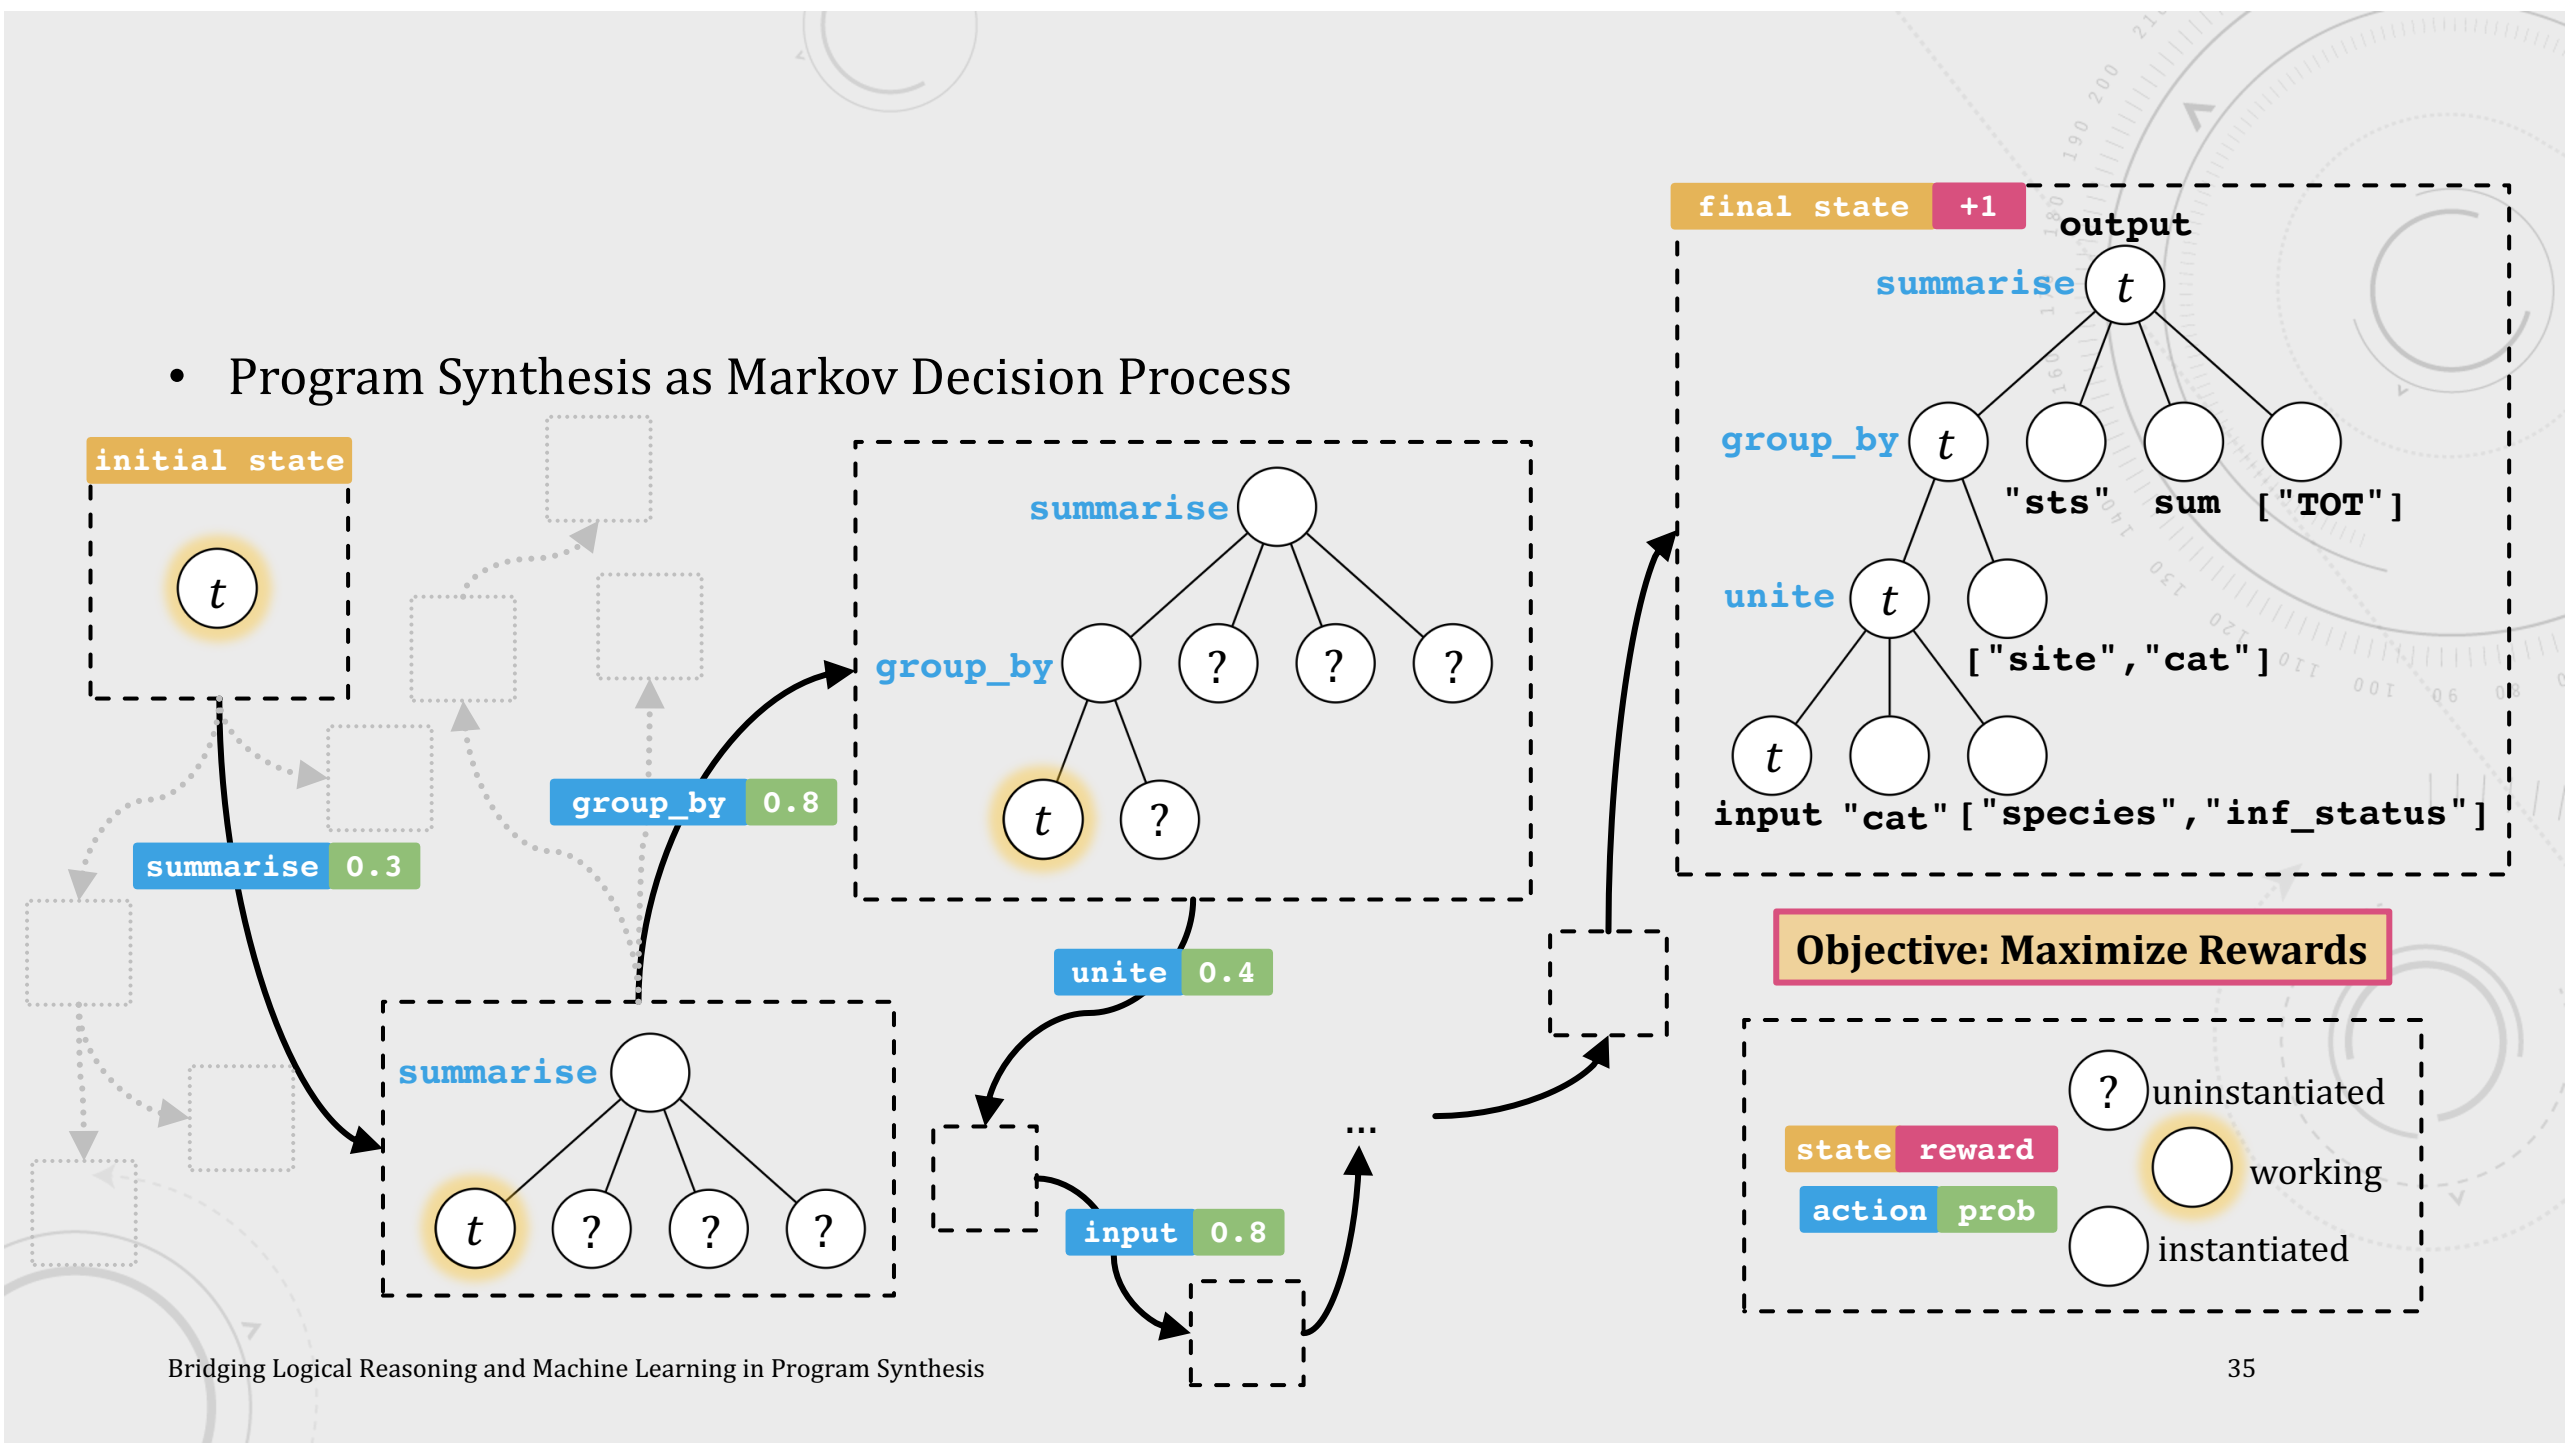
<pre> 1 # gcc -O3 2 3 .L0: 4 movq rsi, r9 5 movl ecx, ecx 6 shrq 32, rsi 7 andl 0xffffffff, r9d 8 movq rcx, rax 9 movl edx, edx 10 imulq r9, rax 11 imulq rdx, r9 12 imulq rsi, rdx 13 imulq rsi, rcx 14 addq rdx, rax 15 jae .L2 16 movabsq 0x100000000, rdx 17 addq rdx, rcx 18 .L2: 19 movq rax, rsi 20 movq rax, rdx 21 shrq 32, rsi 22 salq 32, rdx 23 addq rsi, rcx 24 addq r9, rdx 25 adcq 0, rcx 26 addq r8, rdx 27 adcq 0, rcx 28 addq rdi, rdx 29 adcq 0, rcx 30 movq rcx, r8 31 movq rdx, rdi </pre>	<pre> 1 # STOKe 2 3 .L0: 4 shlq 32, rcx 5 movl edx, edx 6 xorq rdx, rcx 7 movq rcx, rax 8 mulq rsi 9 addq r8, rdi 10 adcq 0, rdx 11 addq rdi, rax 12 adcq 0, rdx 13 movq rdx, r8 14 movq rax, rdi </pre>
--	---

The goal is to find a sequence of x86 instructions that are close to optimal for a particular task

Figure 1. Montgomery multiplication kernel from the OpenSSL big number library, compiled by gcc -O3 (left) and STOKe (right). The STOKe code is 16 lines shorter, 1.6x faster, and slightly faster than expert handwritten assembly.

Markov Chains

- Program Synthesis as Markov Decision Process



Cost function

$$c(\mathcal{R}; \mathcal{T}) = \text{eq}(\mathcal{R}; \mathcal{T}) + \text{perf}(\mathcal{R}; \mathcal{T}) \quad (2)$$

$$\mathcal{R}' = \arg \min_r \left(\text{perf}(r; \mathcal{T}) \mid \text{eq}(r; \mathcal{T}) = 0 \right) \quad (3)$$


The transformation correctness term, $\text{eq}(\cdot)$, measures the similarity of two functions. The term is zero if and only if the two functions are equal. The performance improvement term, $\text{perf}(\cdot)$, quantifies the performance improvement of a rewrite with respect to the target. Depending on the application, this term could reflect code size, expected runtime, number of disk accesses, power consumption, or any other measure of resource usage.

Markov chain Monte Carlo (MCMC) sampling

- Select an initial program
- Repeat (millions to billions of times)
 - Propose a random change and evaluate cost
 - If (decreased) {accept}
 - If (increased) {may accept}

MCMC sampling


- Guarantees: Draw samples in proportion to their value; higher value points are sampled more frequently



MCMC sampling

- Intelligent “hill climbing” method, robust against local minima, desirable limiting properties (Metropolis–Hastings)
- Nothing prevents the use of other search methods. Offers a tradeoff between performance and complexity

Results



Average speedup over llvm -O0 for benchmark kernels

TODOs by next lecture

- Proposal will be due on Wed
- Start to work on your final report/project! (20%)